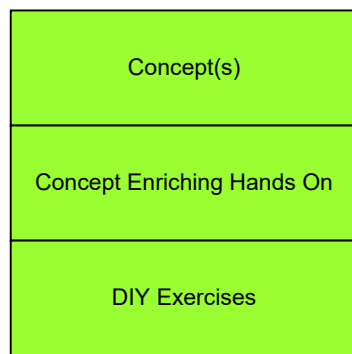
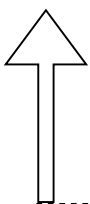
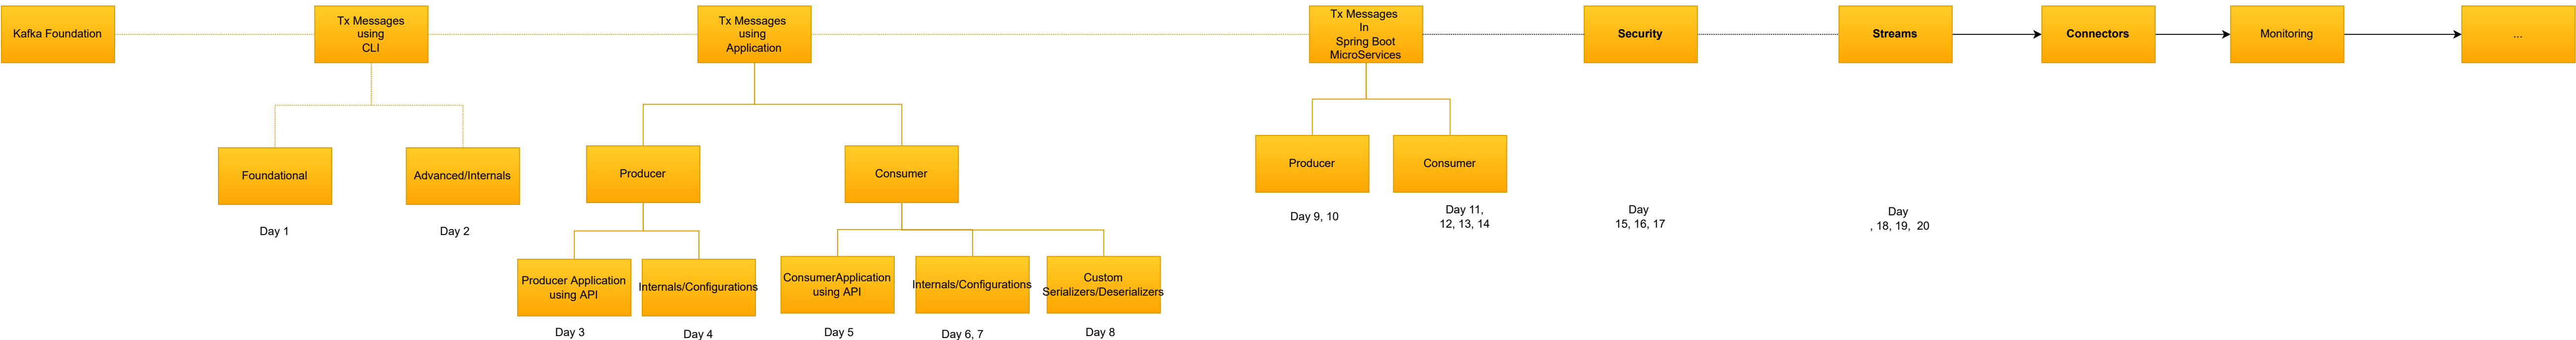
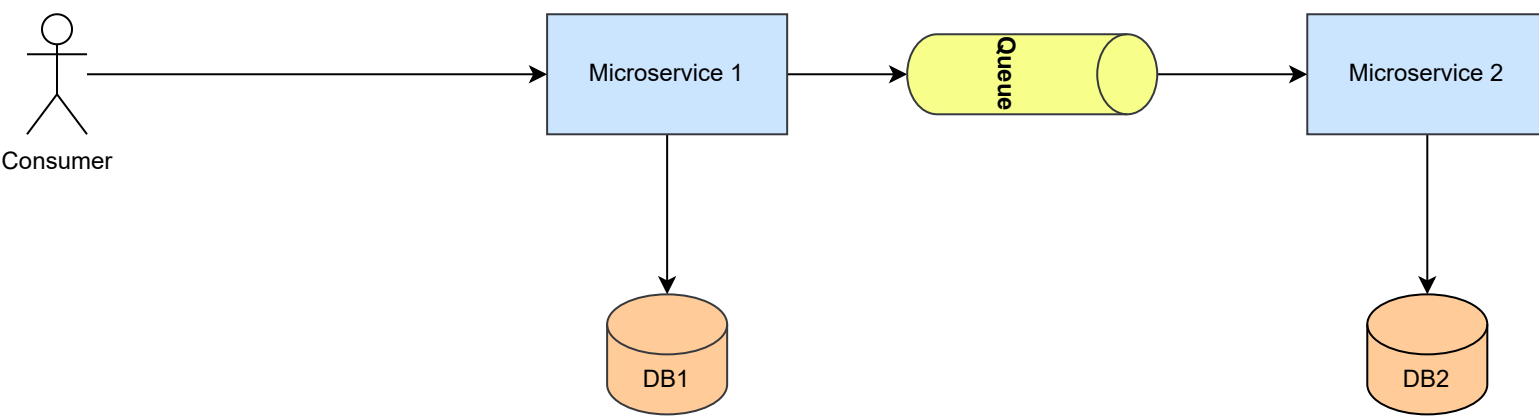
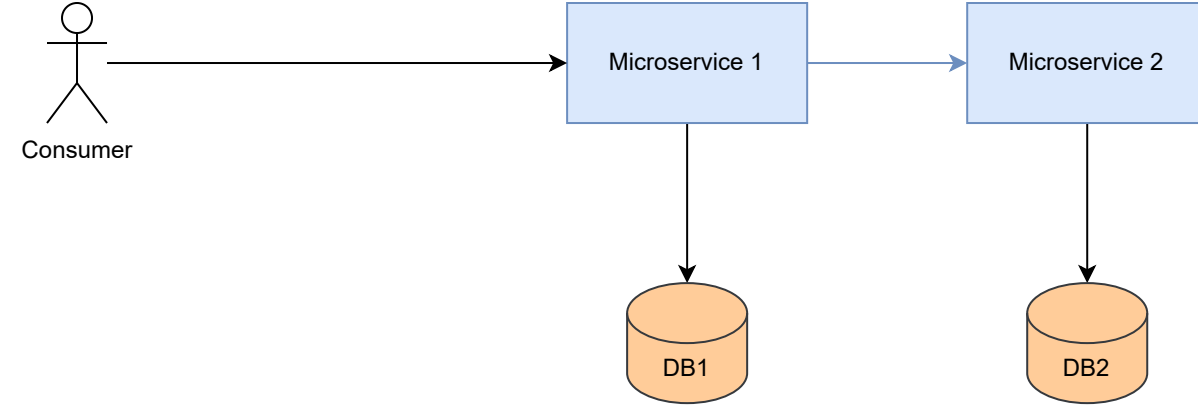


Teaching Principle

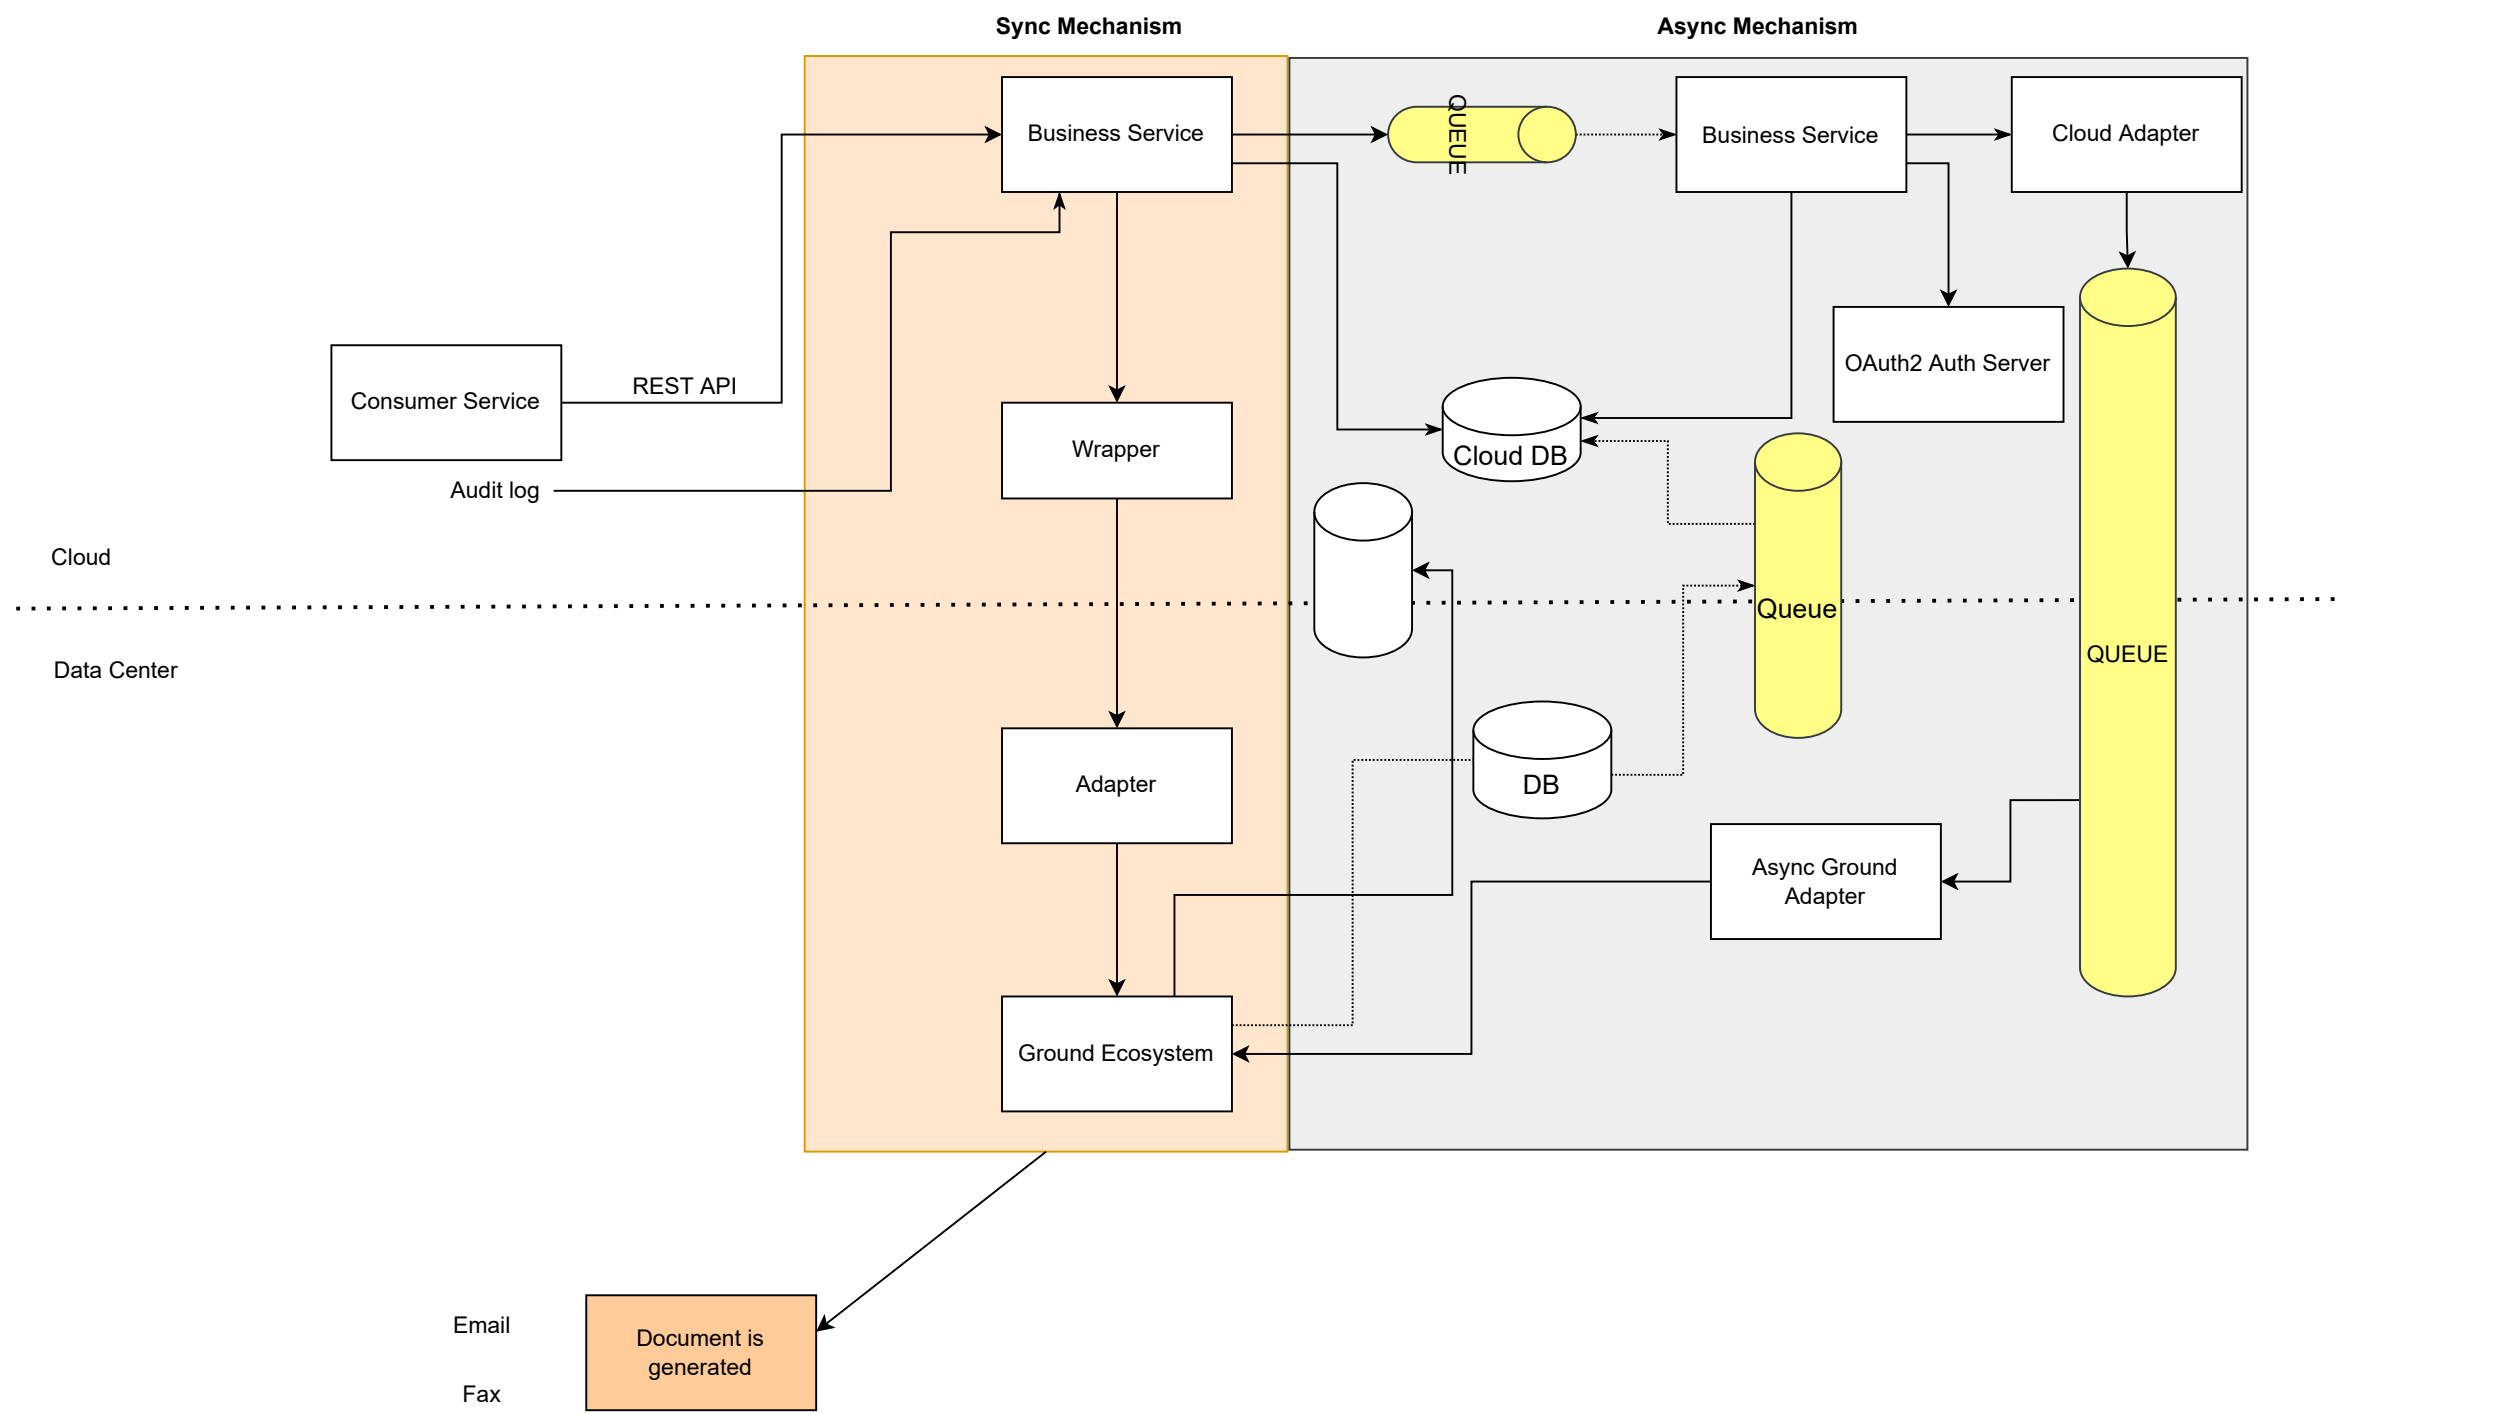


Learning Map

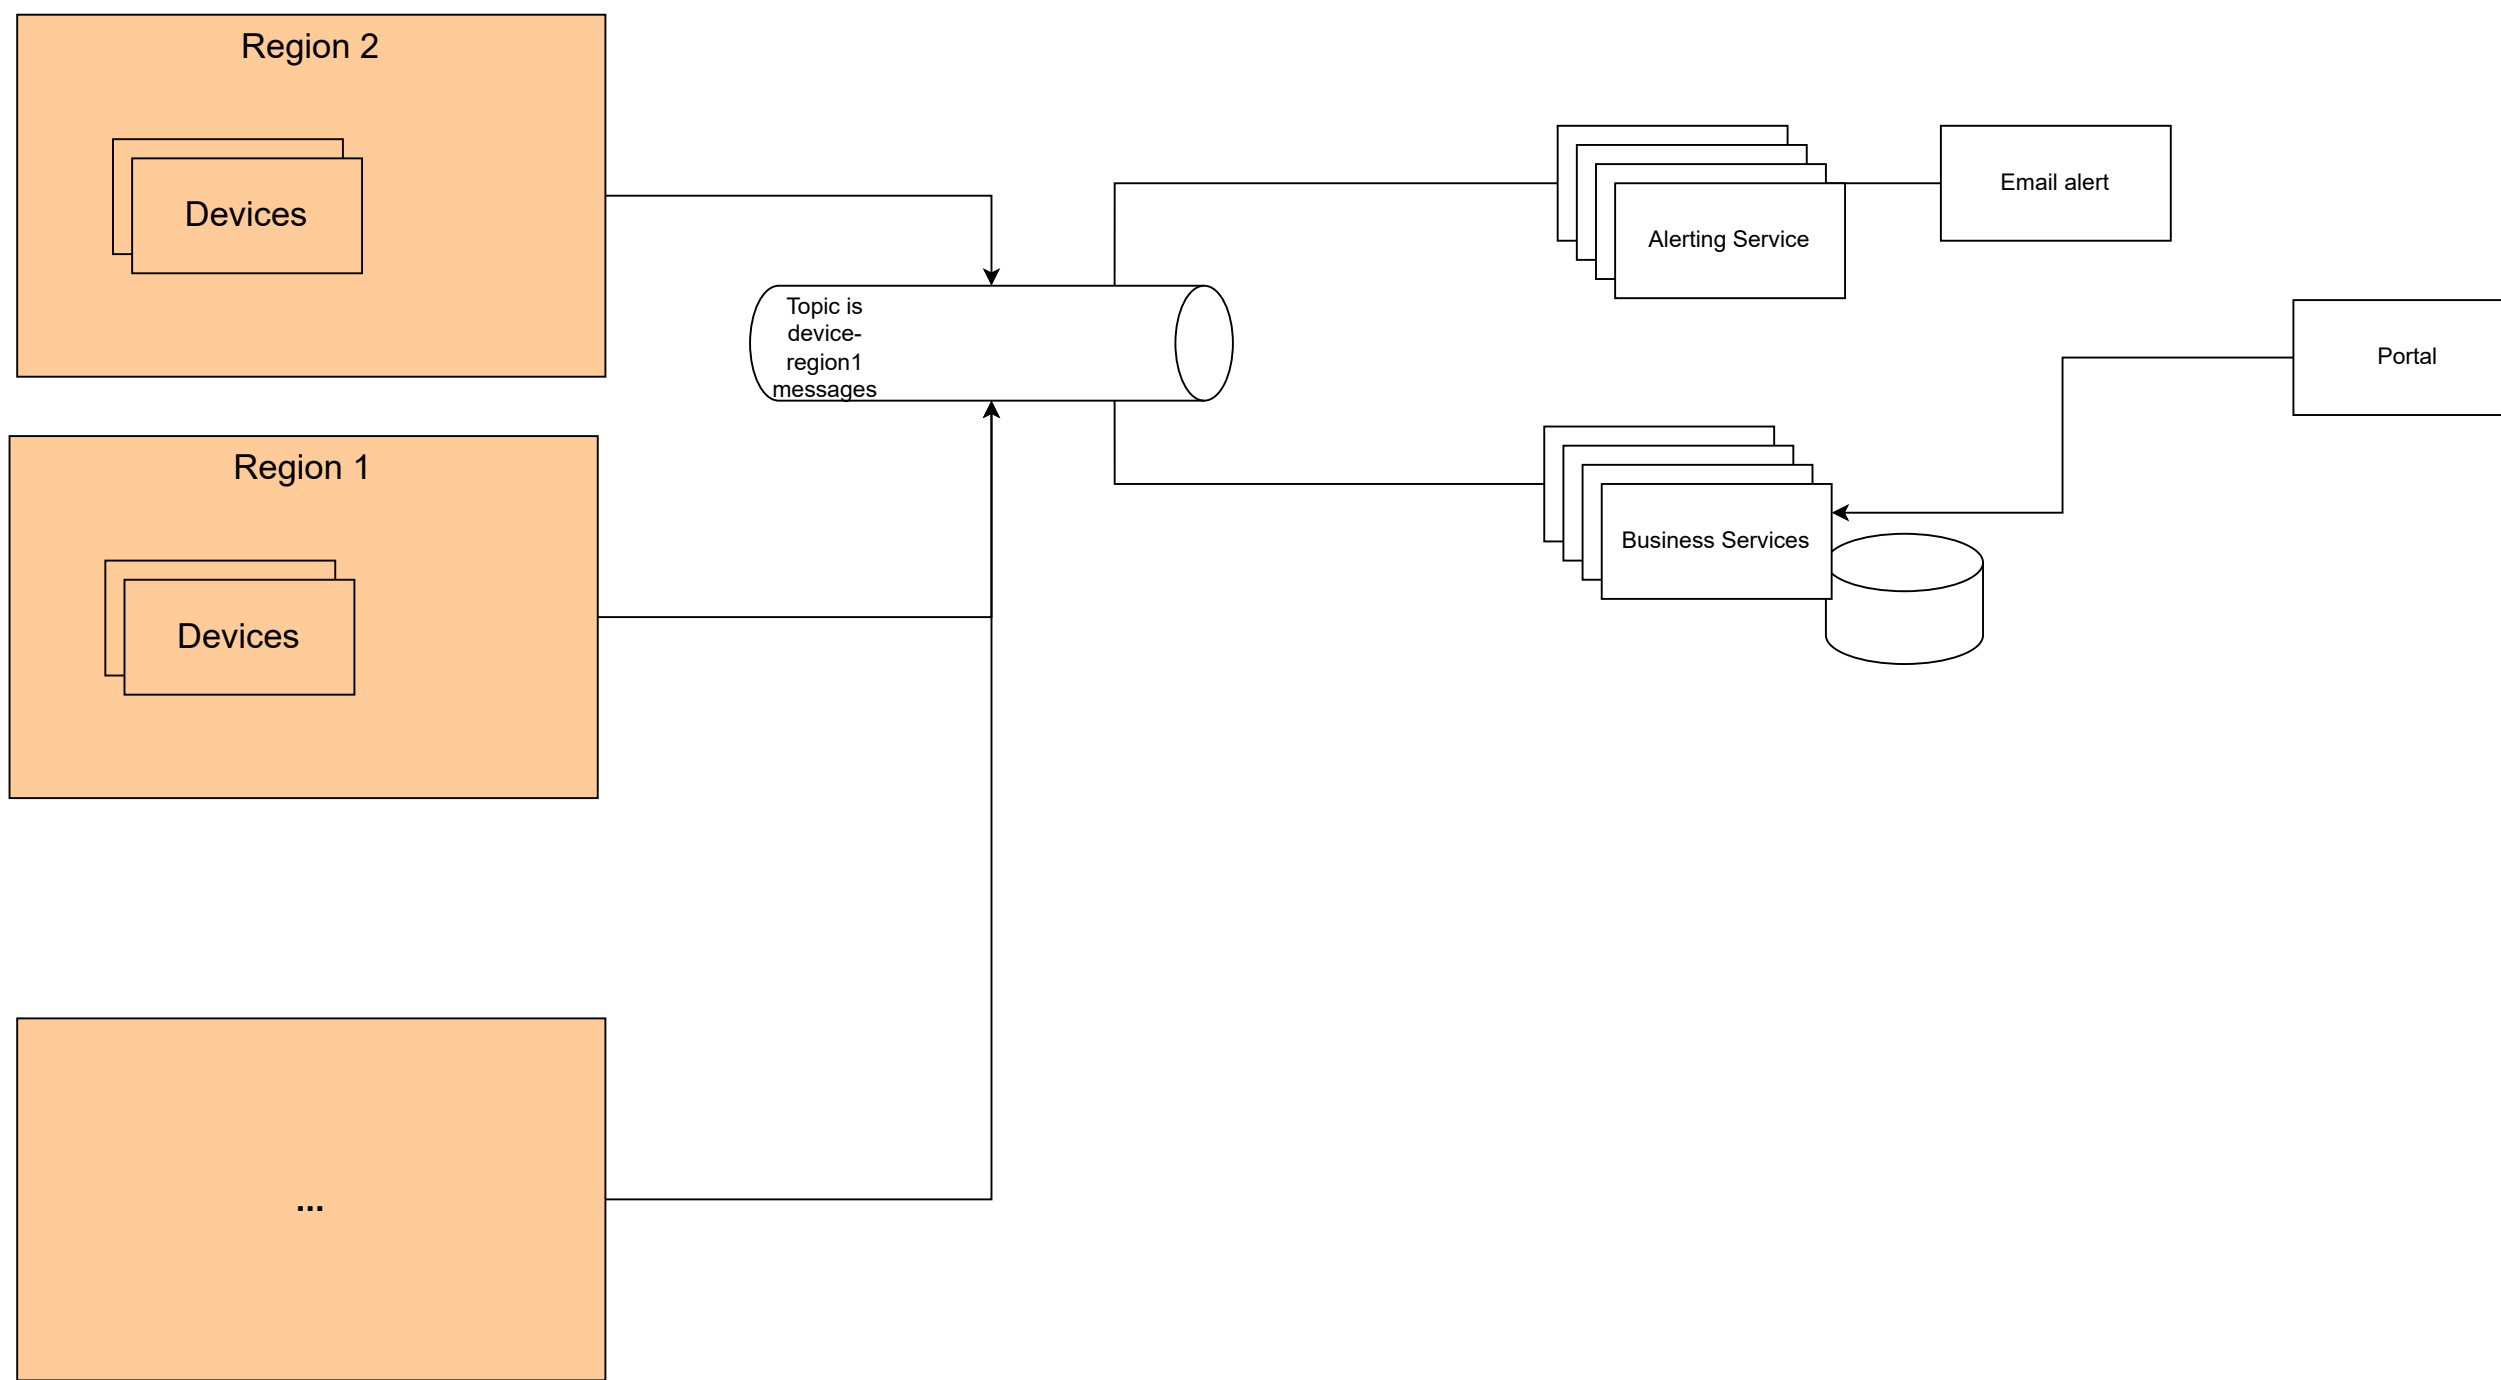




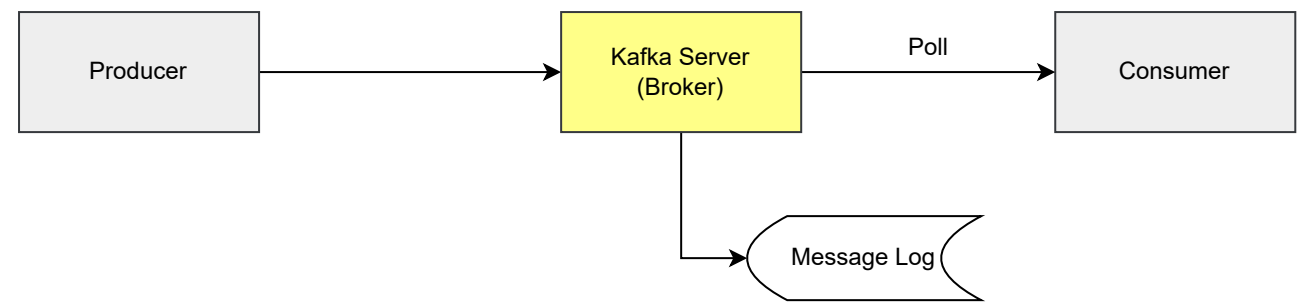
Use case 1



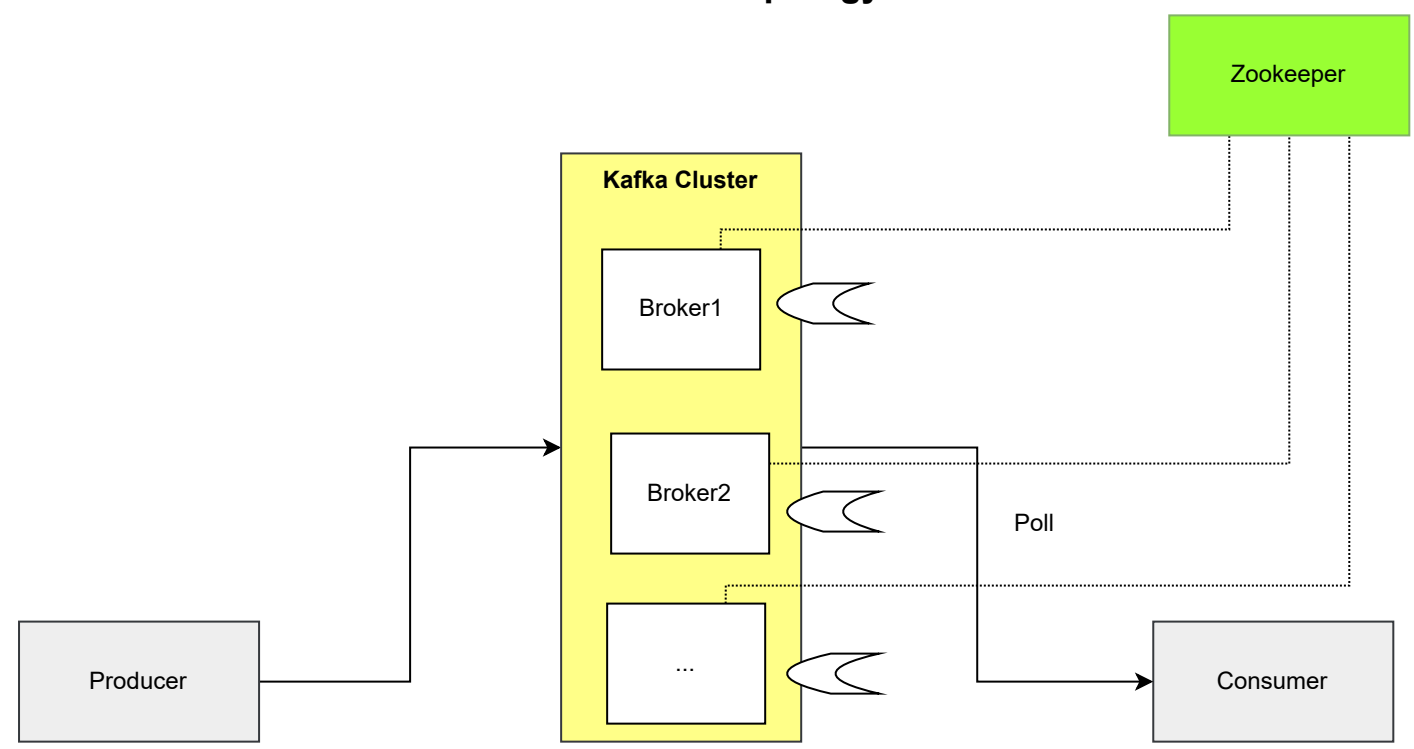
Use case 2



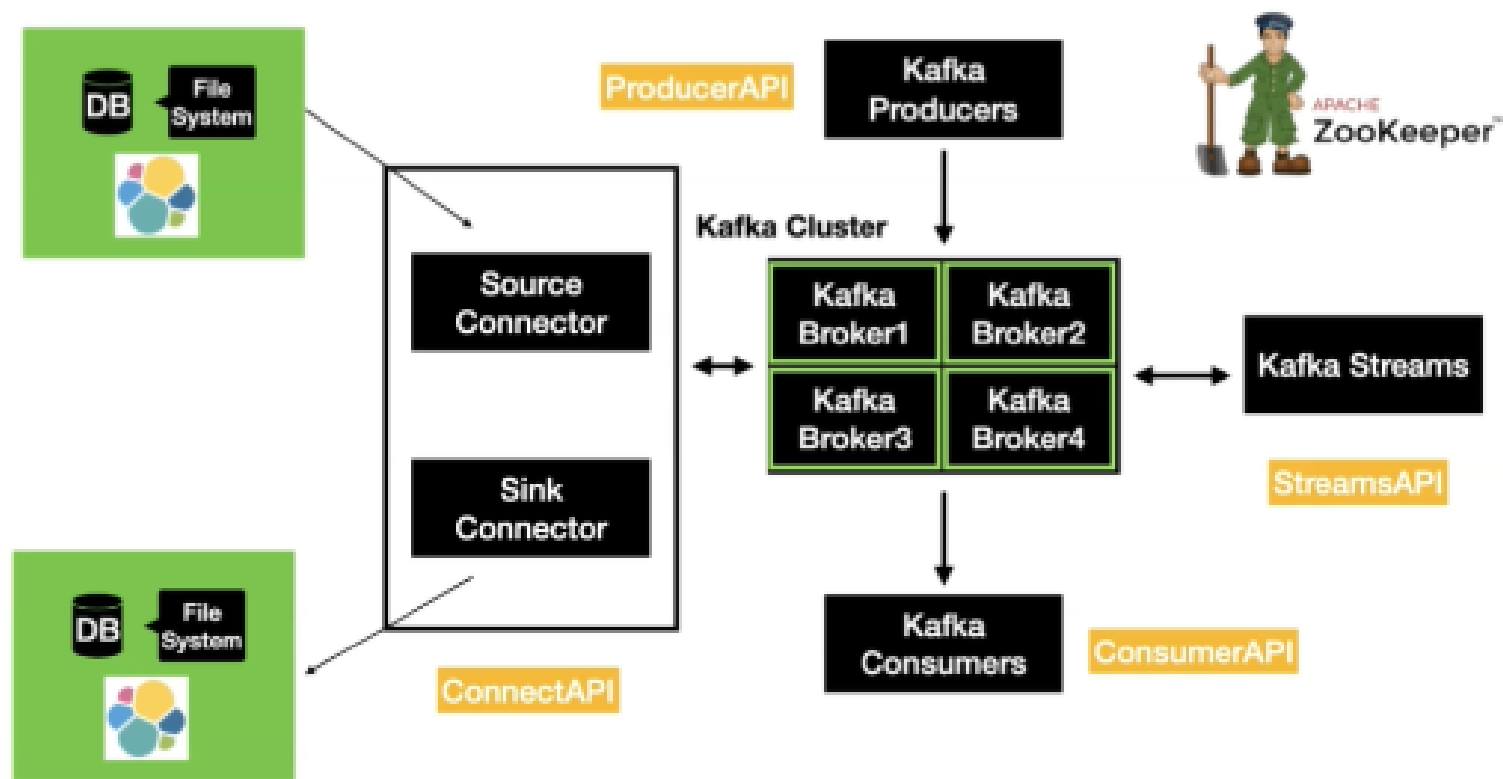
Basic Topology



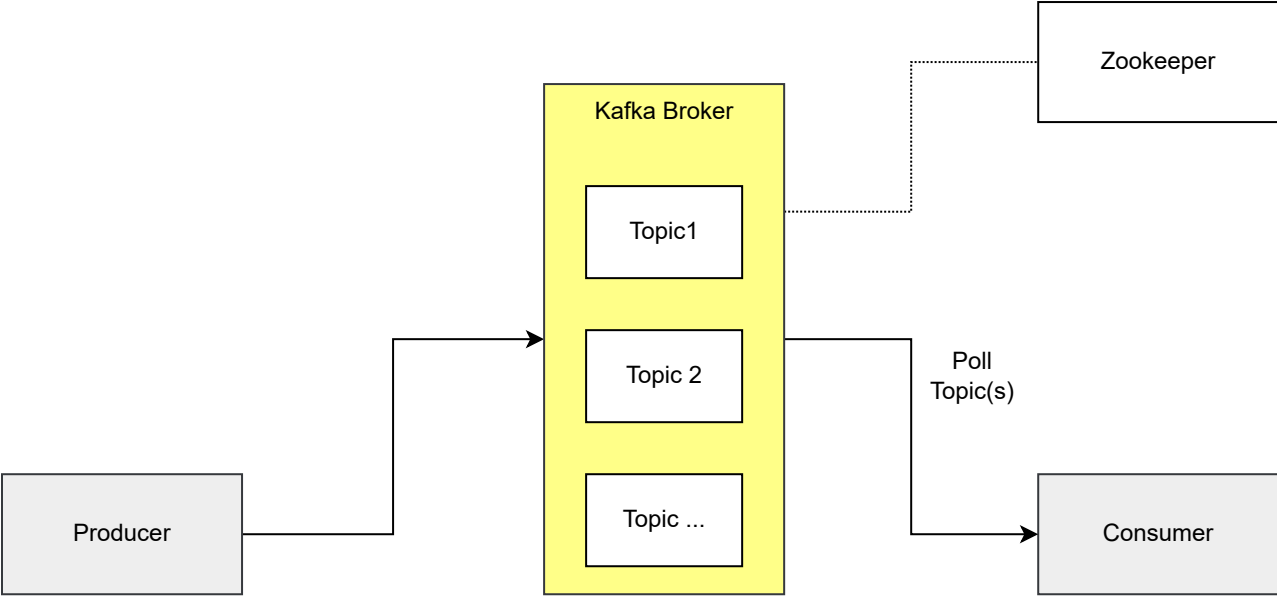
Distributed Topology



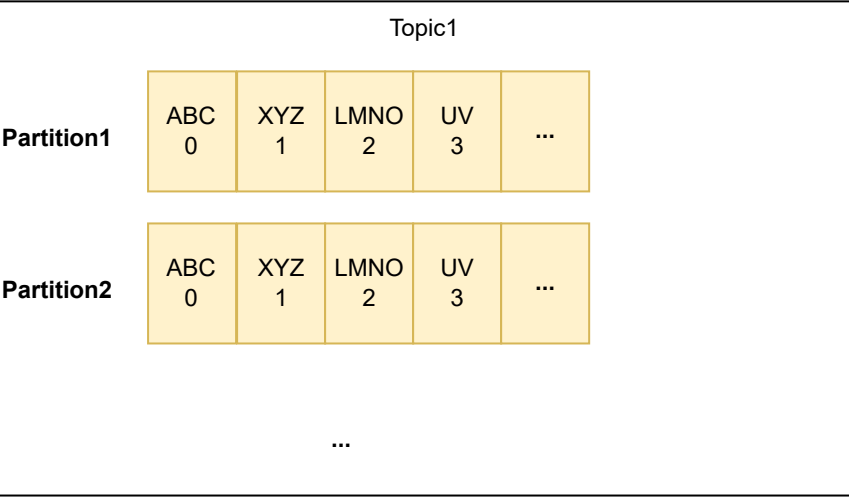
Kafka Landscape



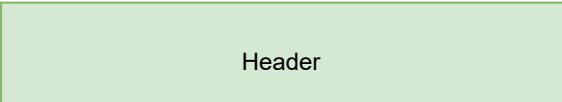
Broker Anatomy



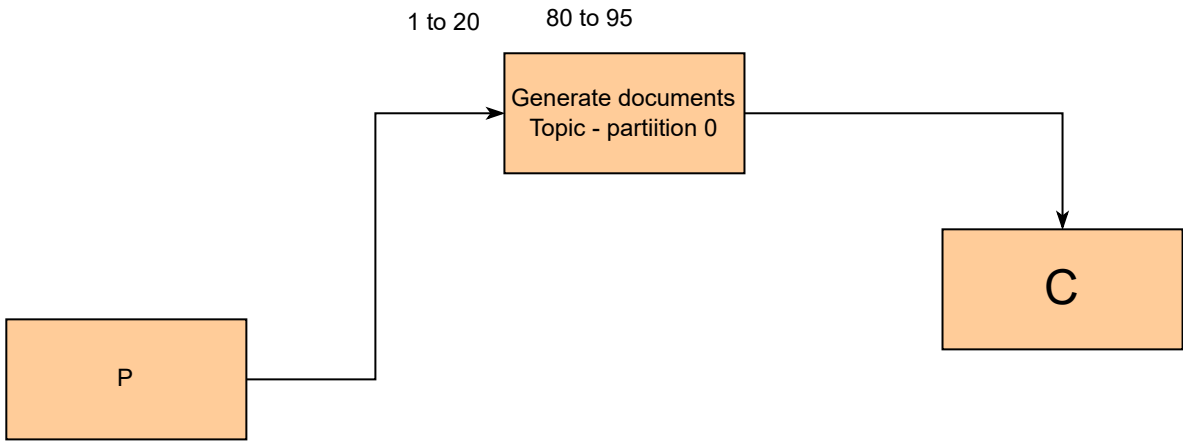
Topic Anatomy



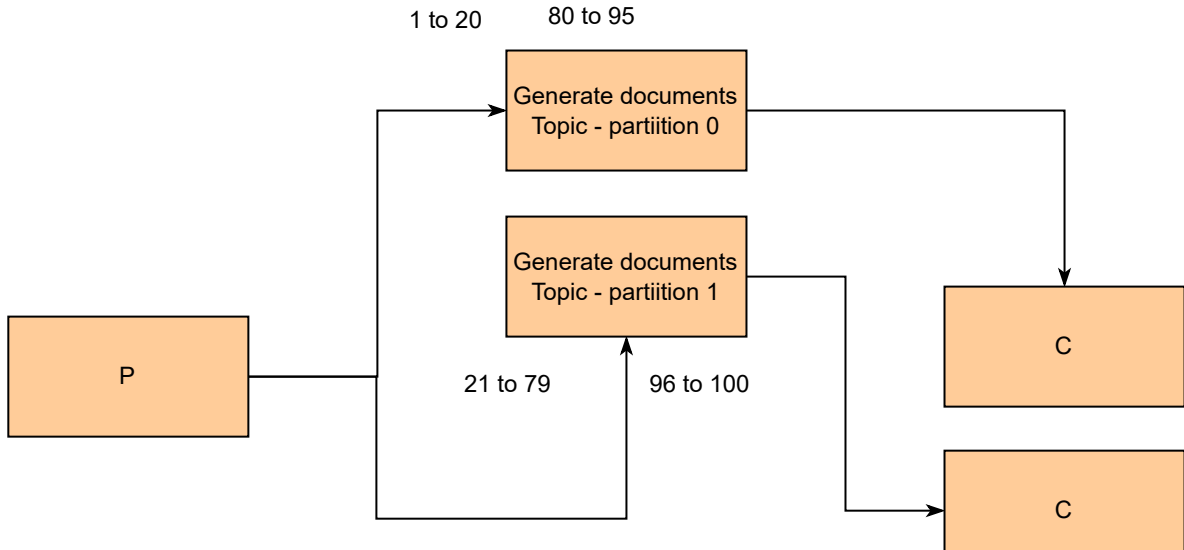
Message Anatomy



Topic-Same



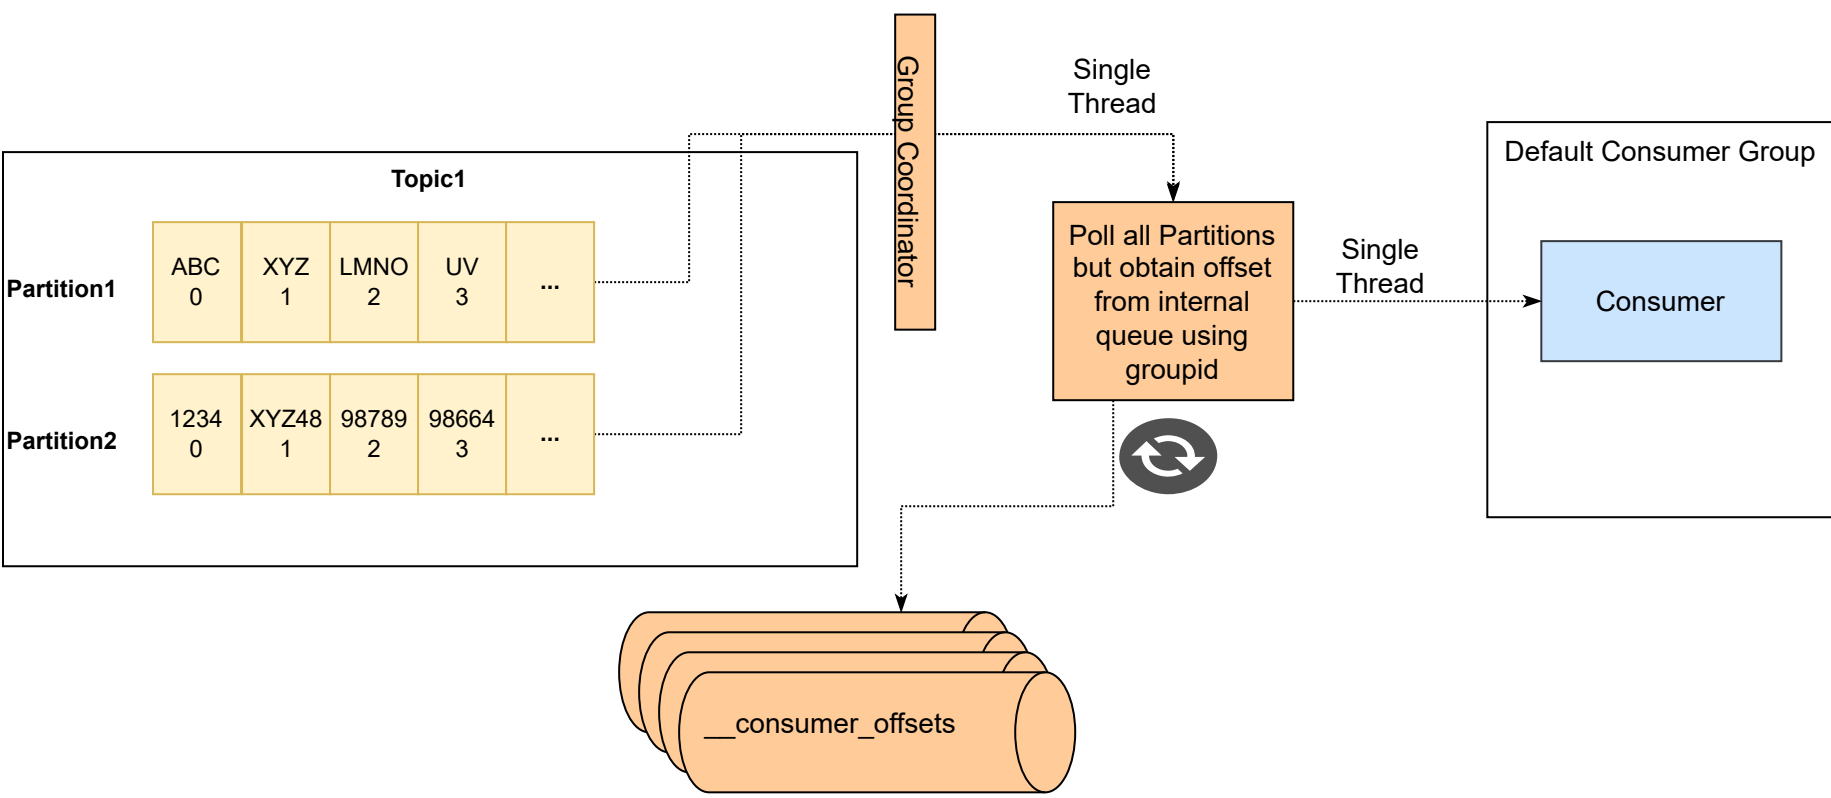
Topic-Same



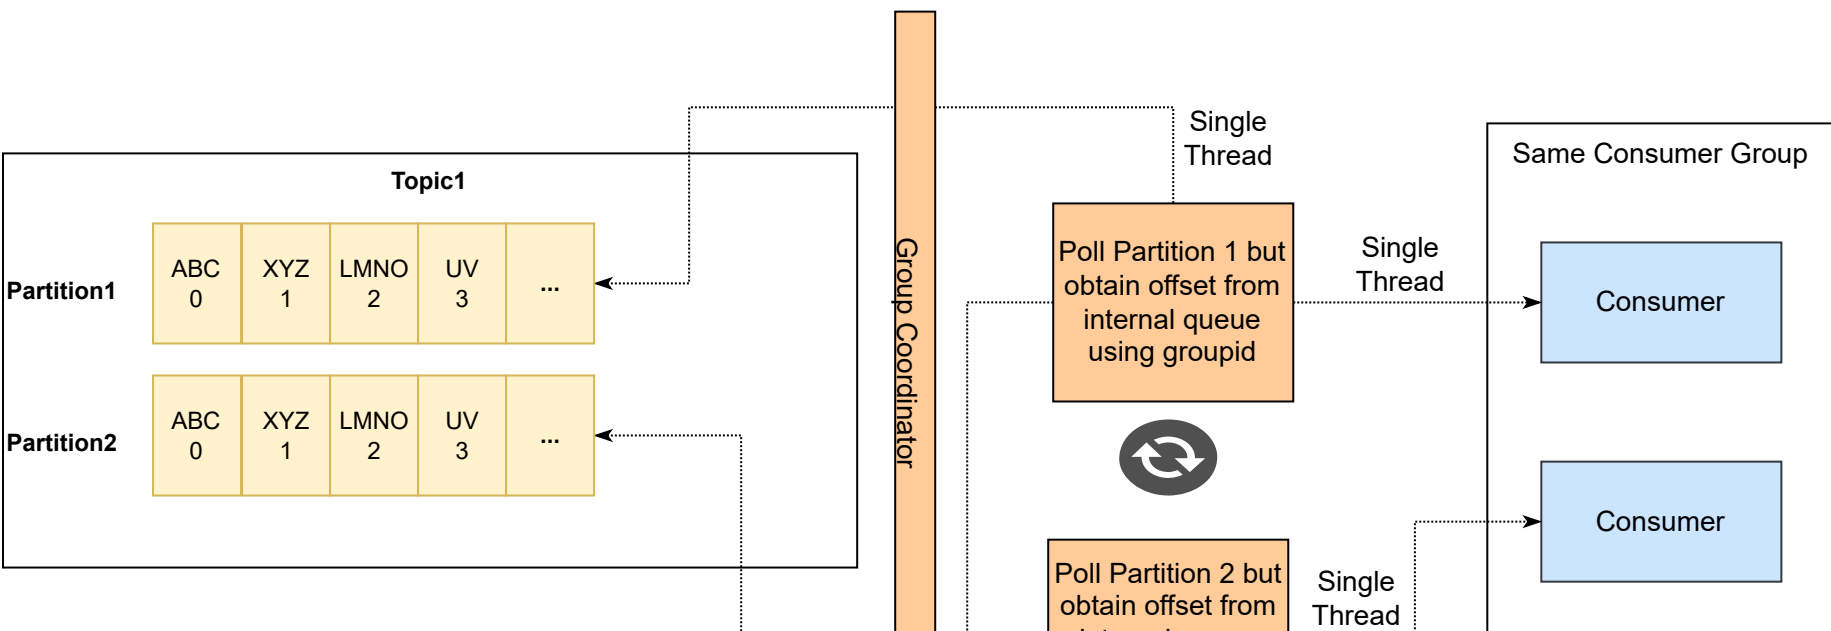
Key(optional)
Value
Timestamp

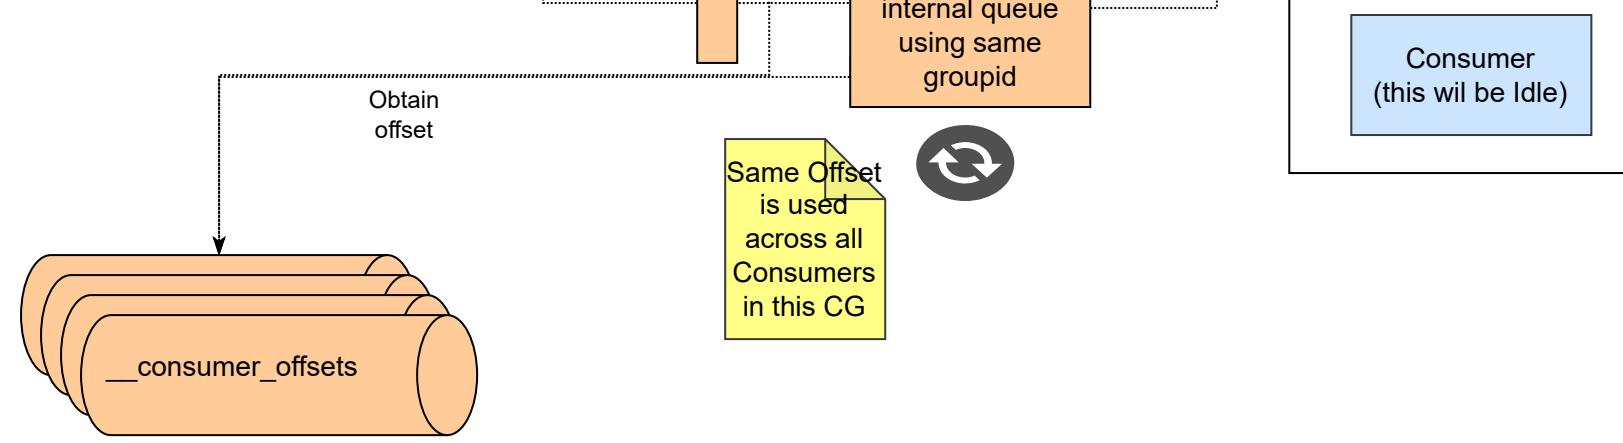


Default Consumer Group

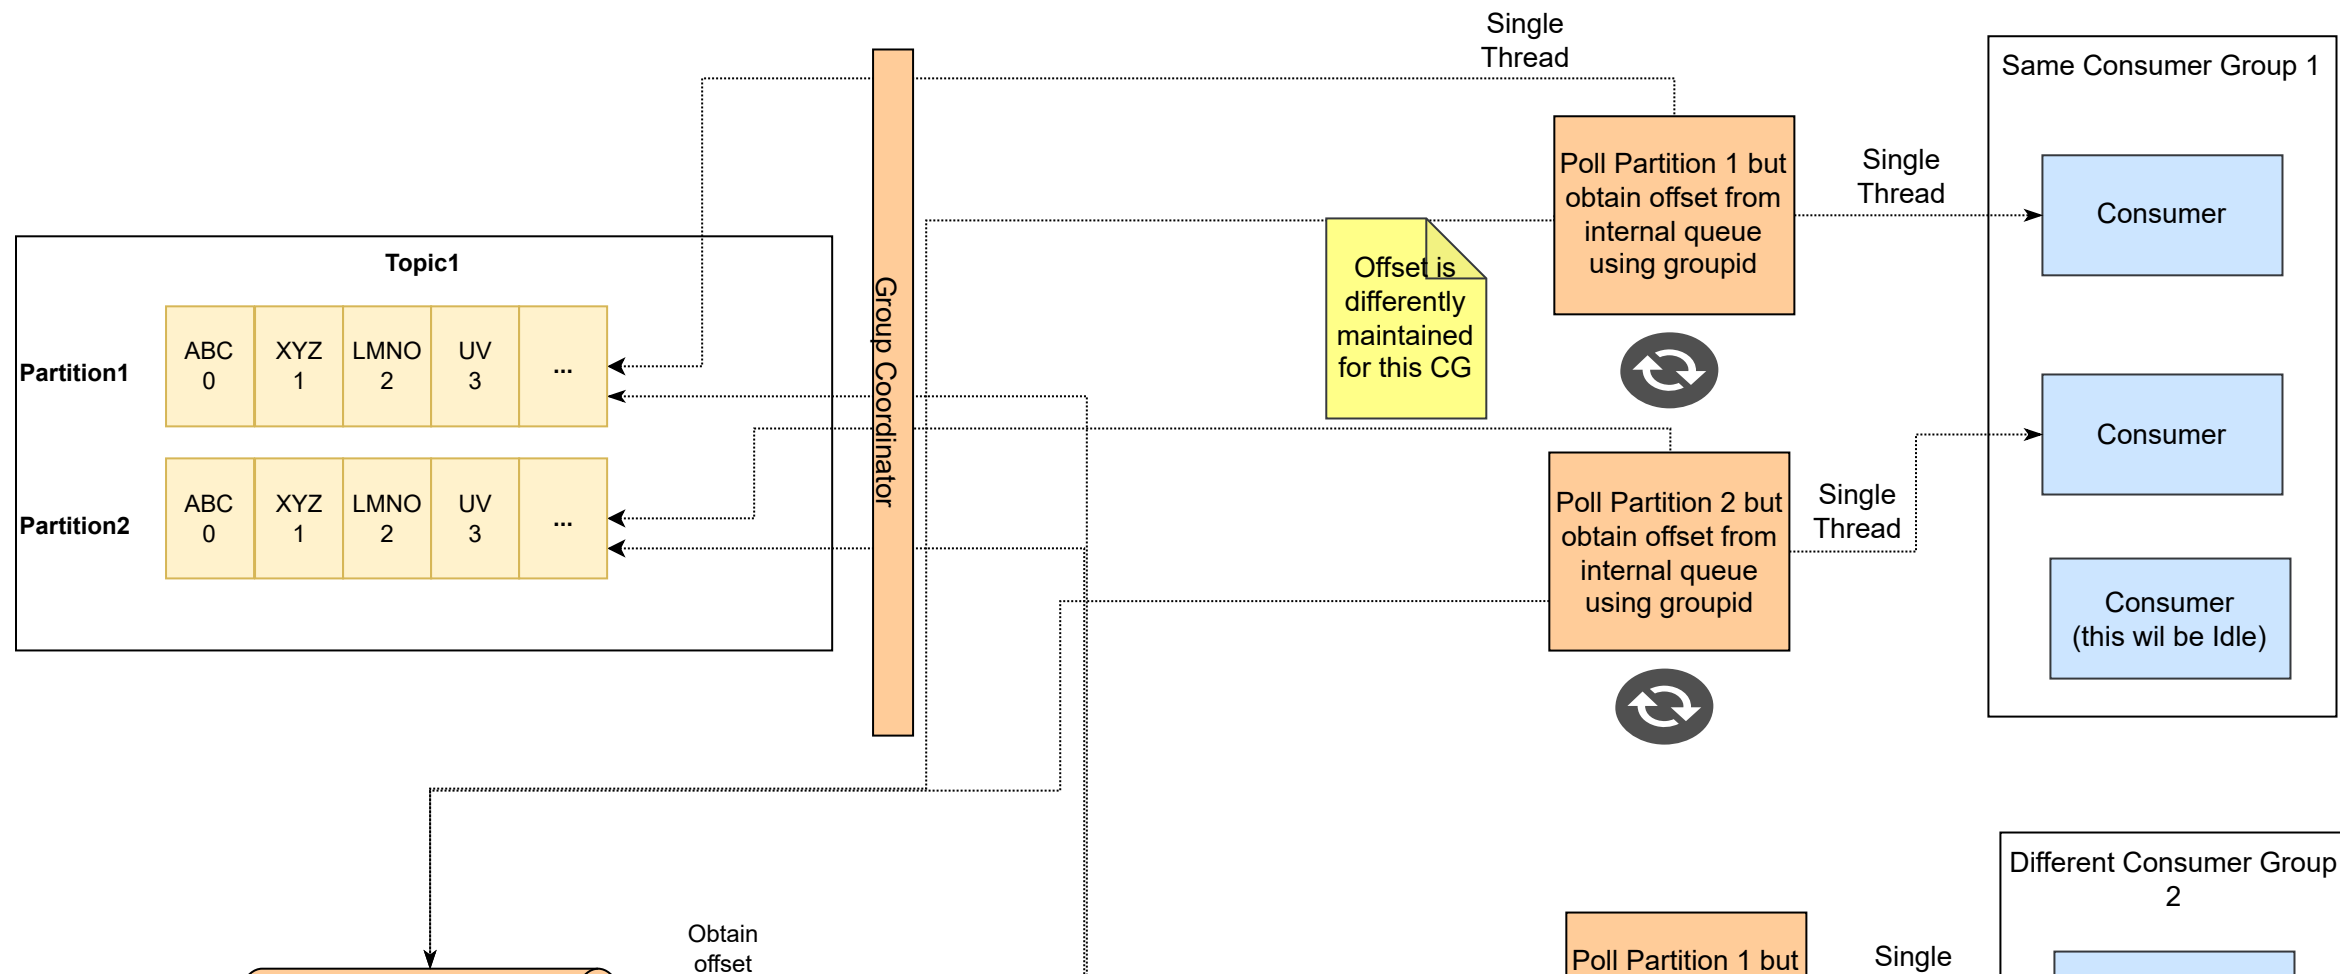


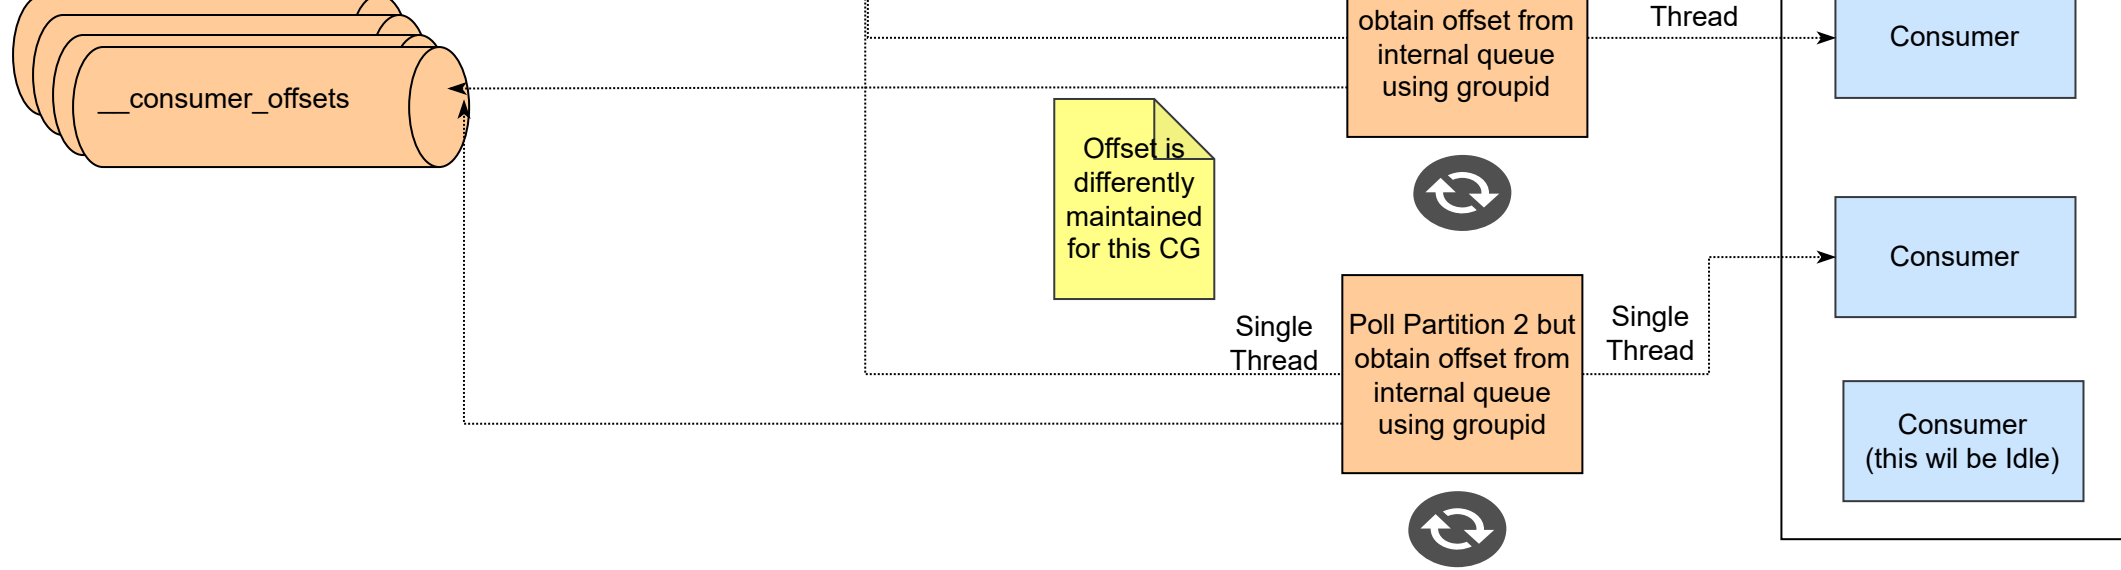
Using same Consumer Group to Read Messages at Scale



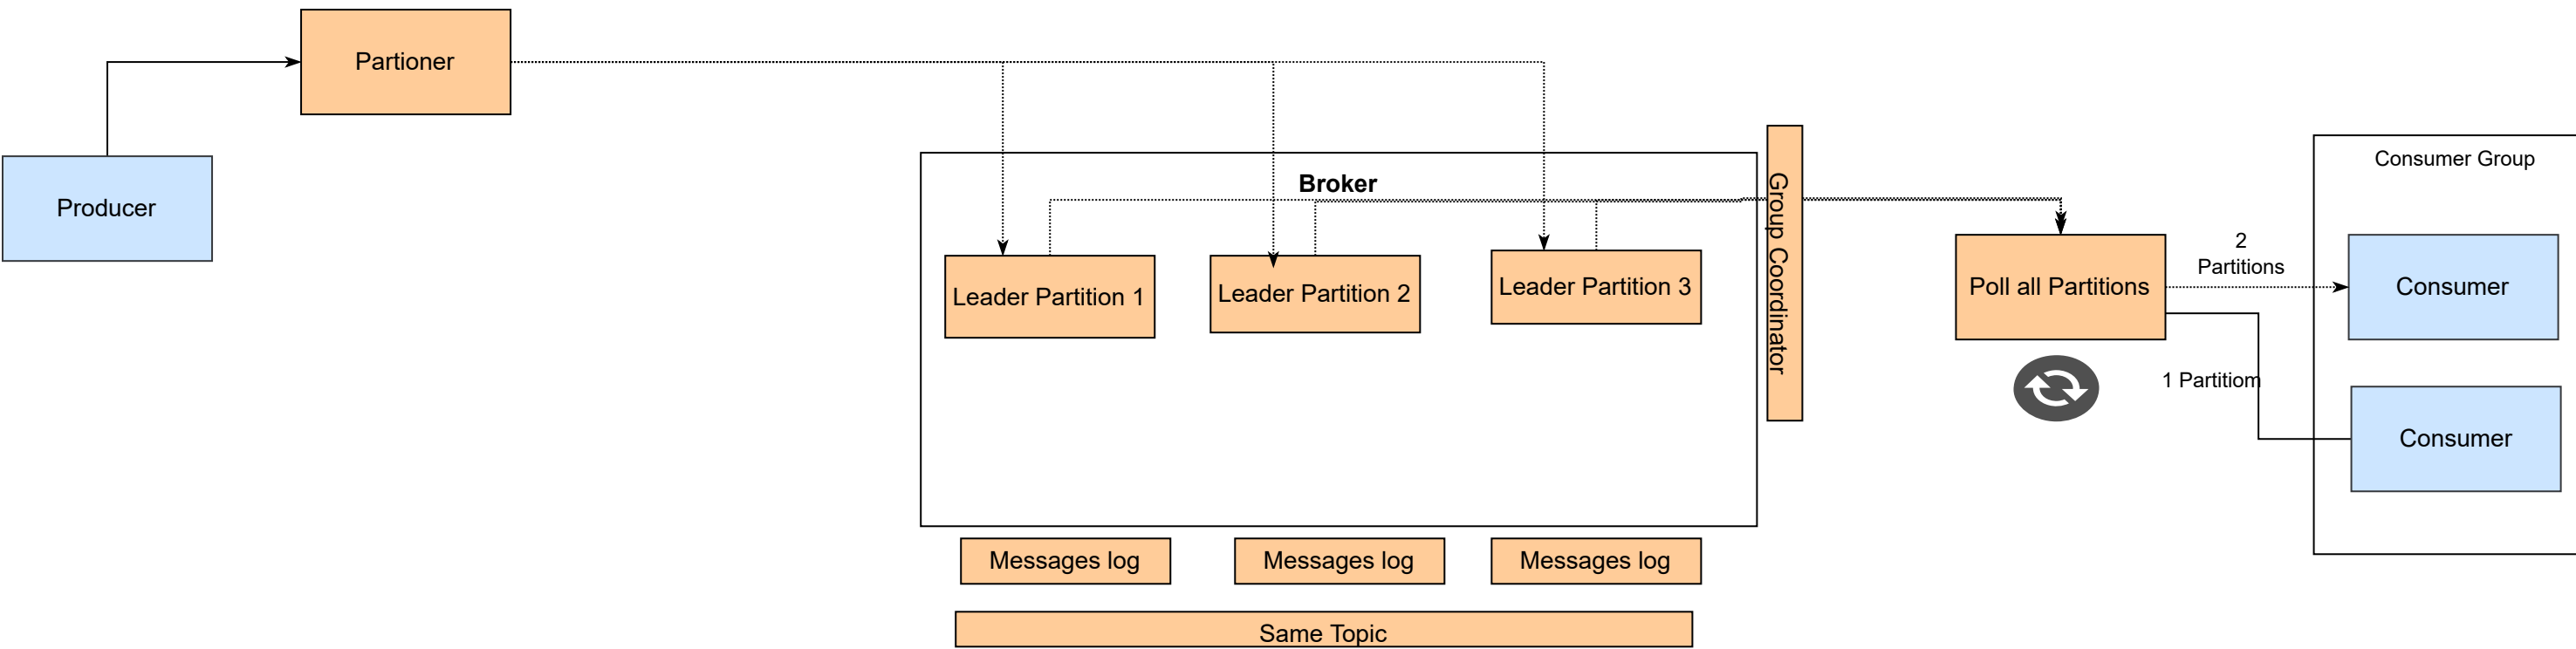


Using Different Consumer Group to Read Messages From Same Topic at Scale

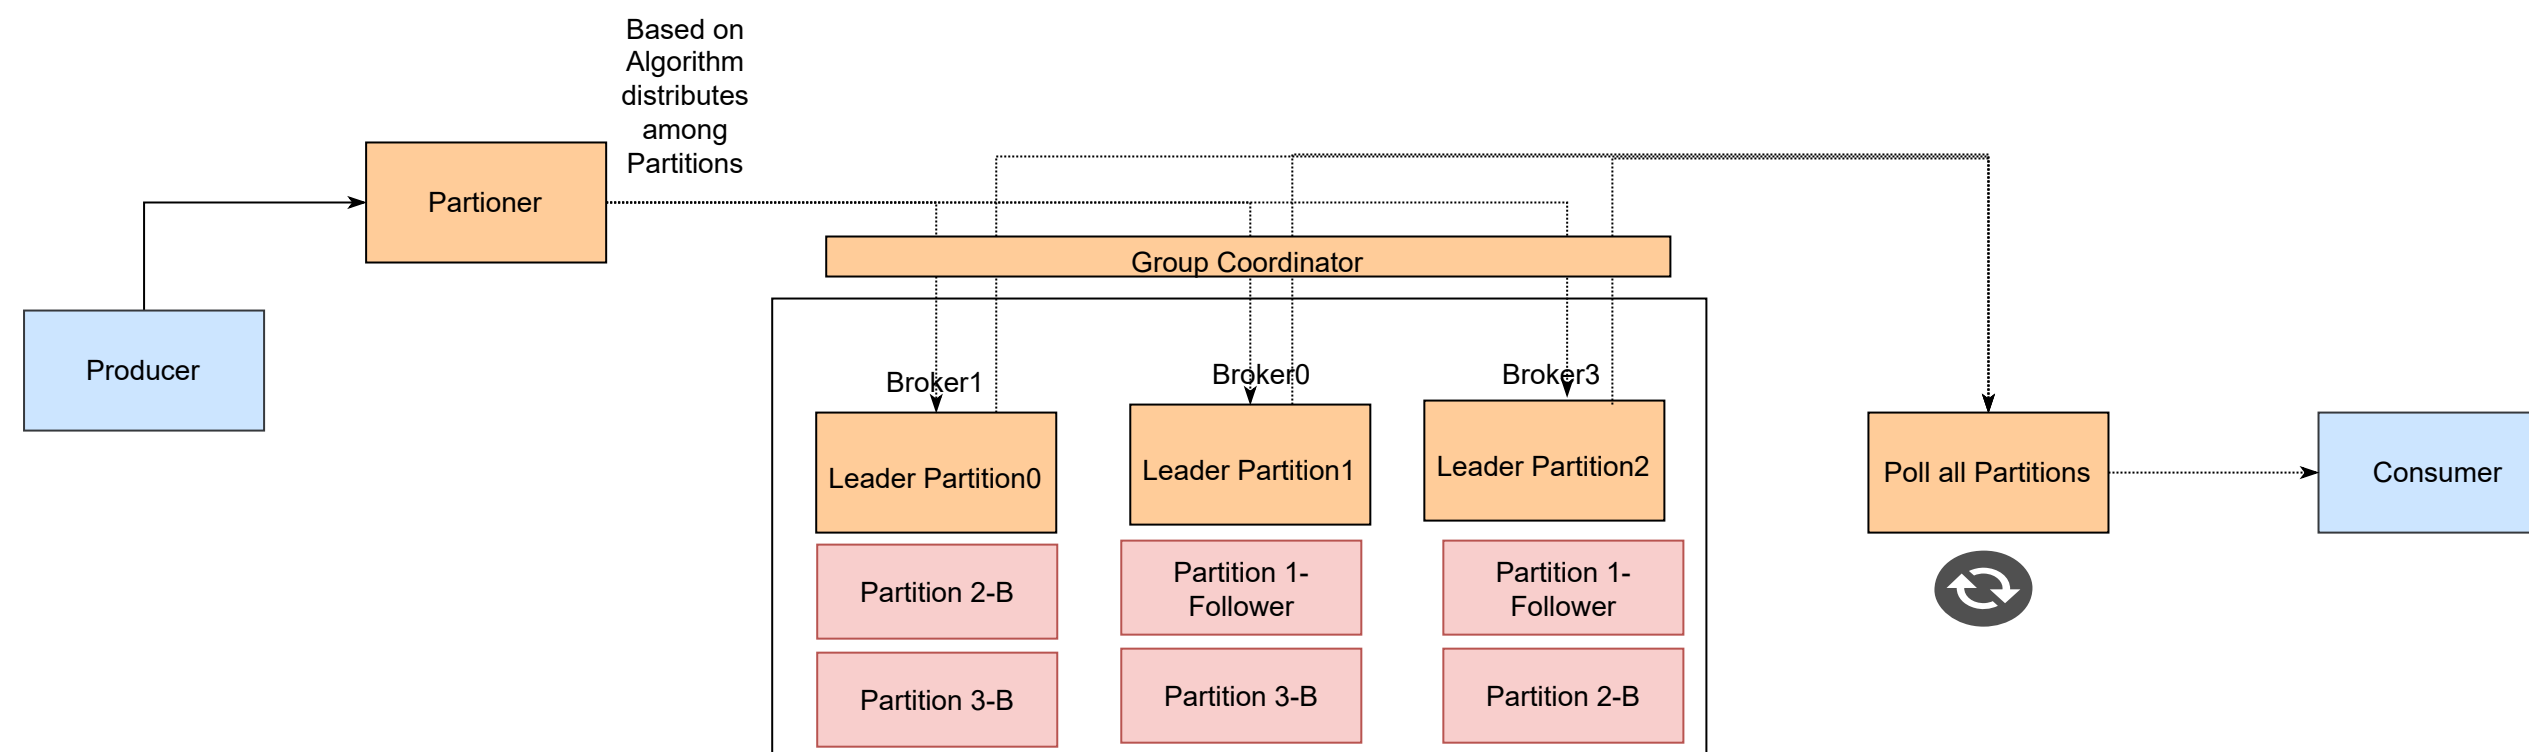




Standalone Topology



Distributed Topology



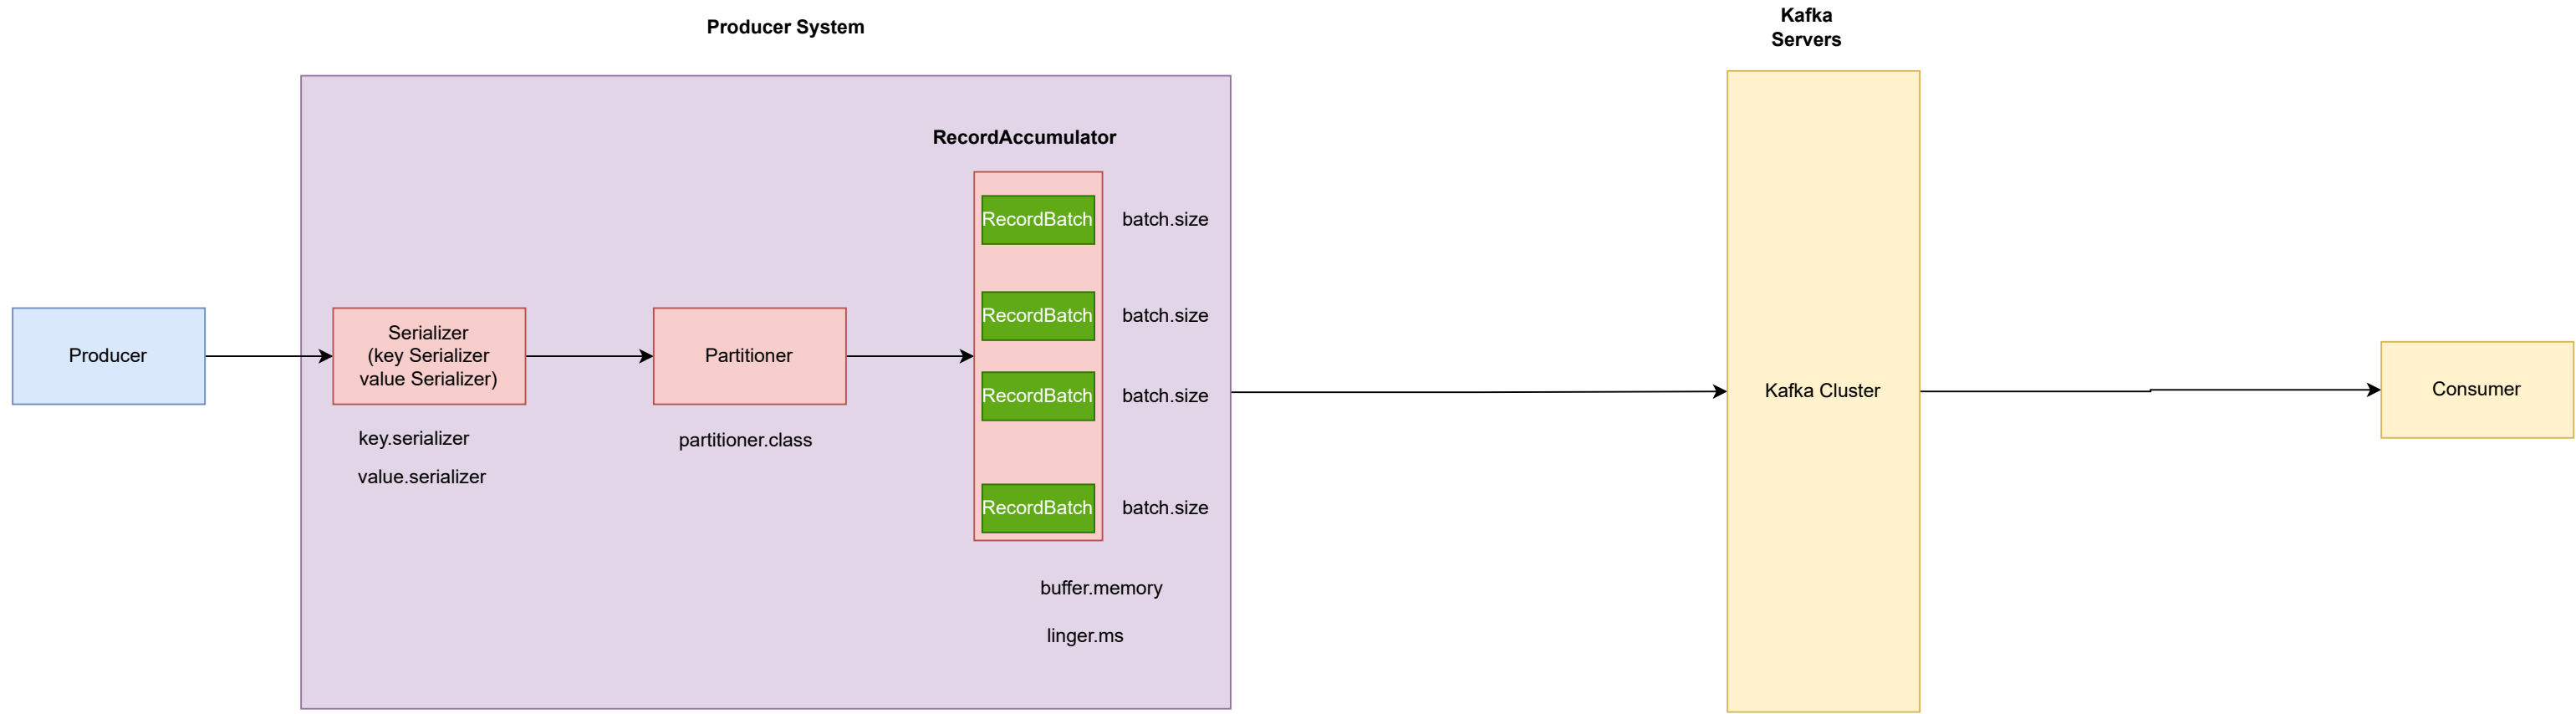
Messages log

Messages log

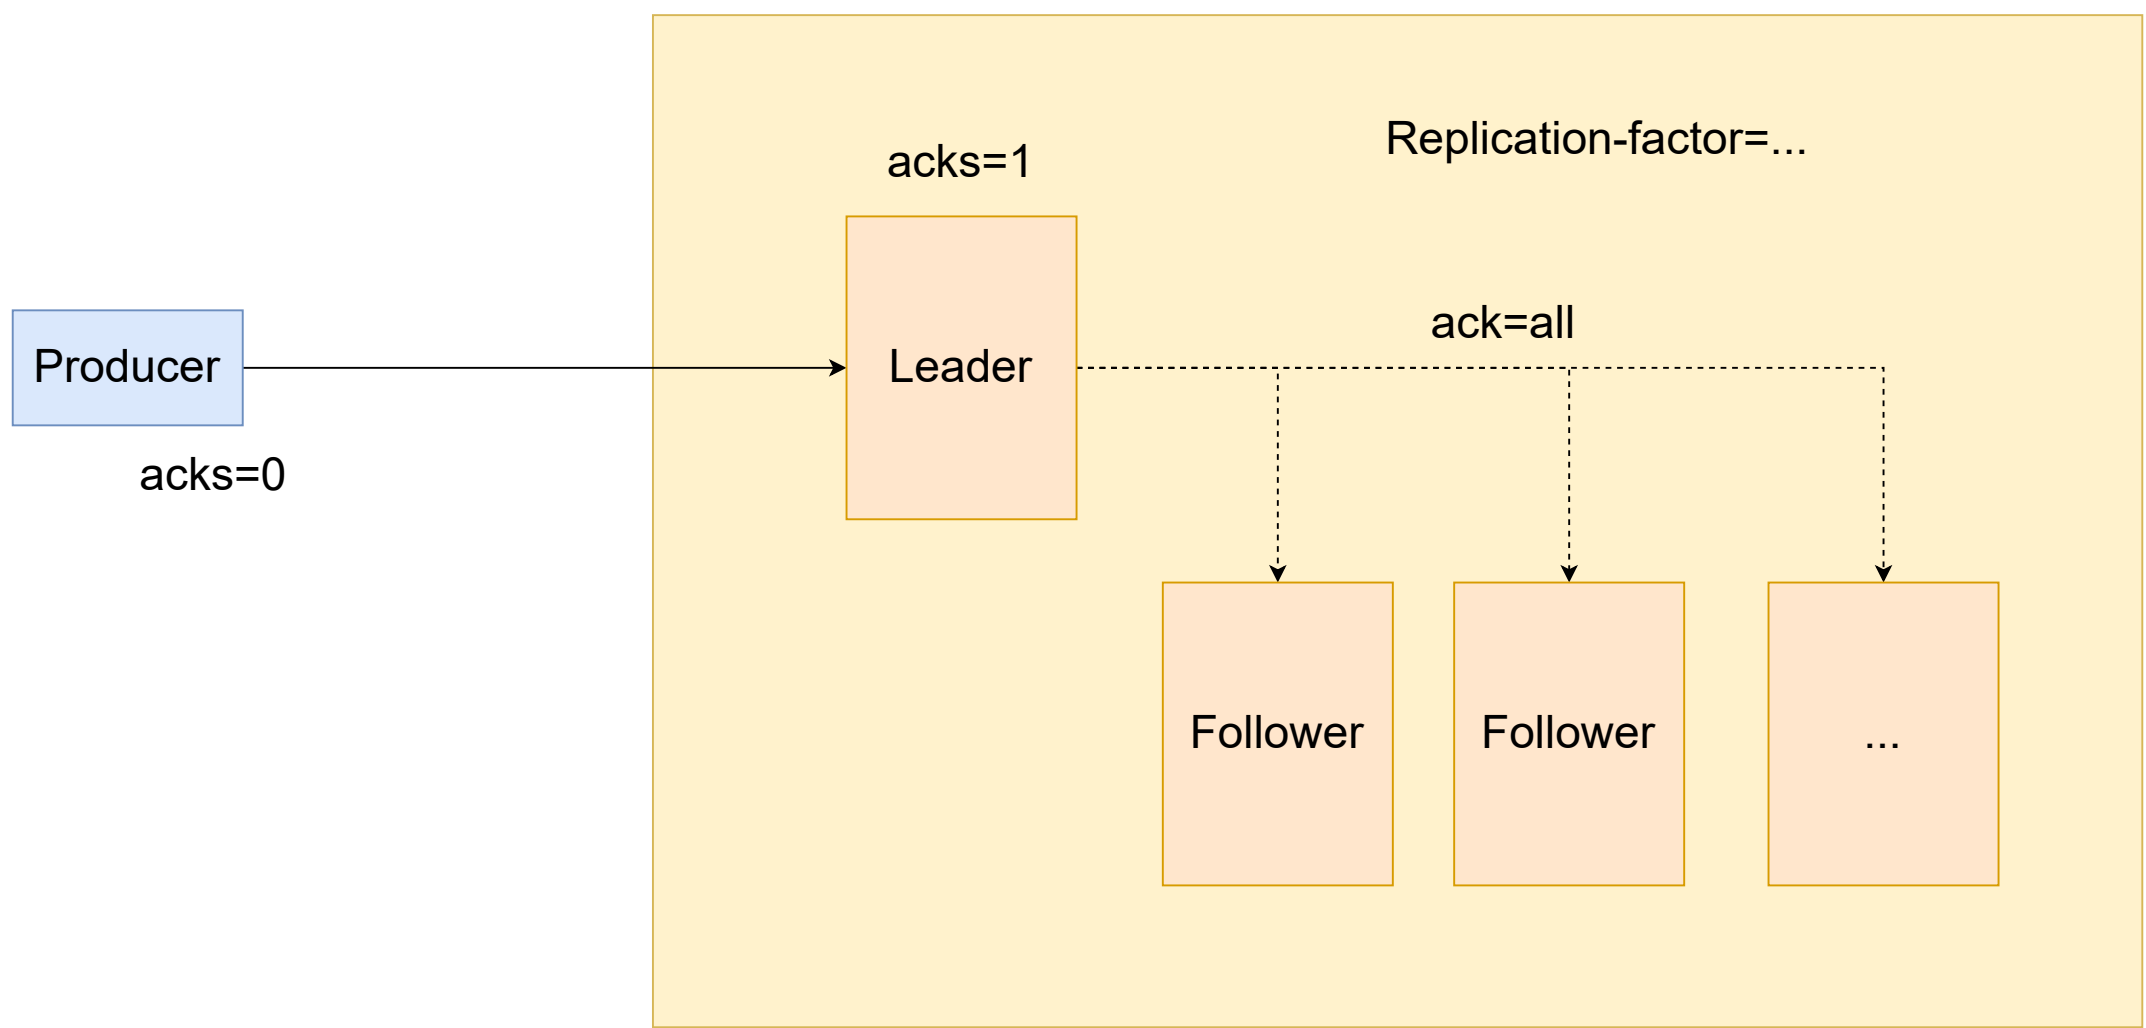
Messages log

Same Topic

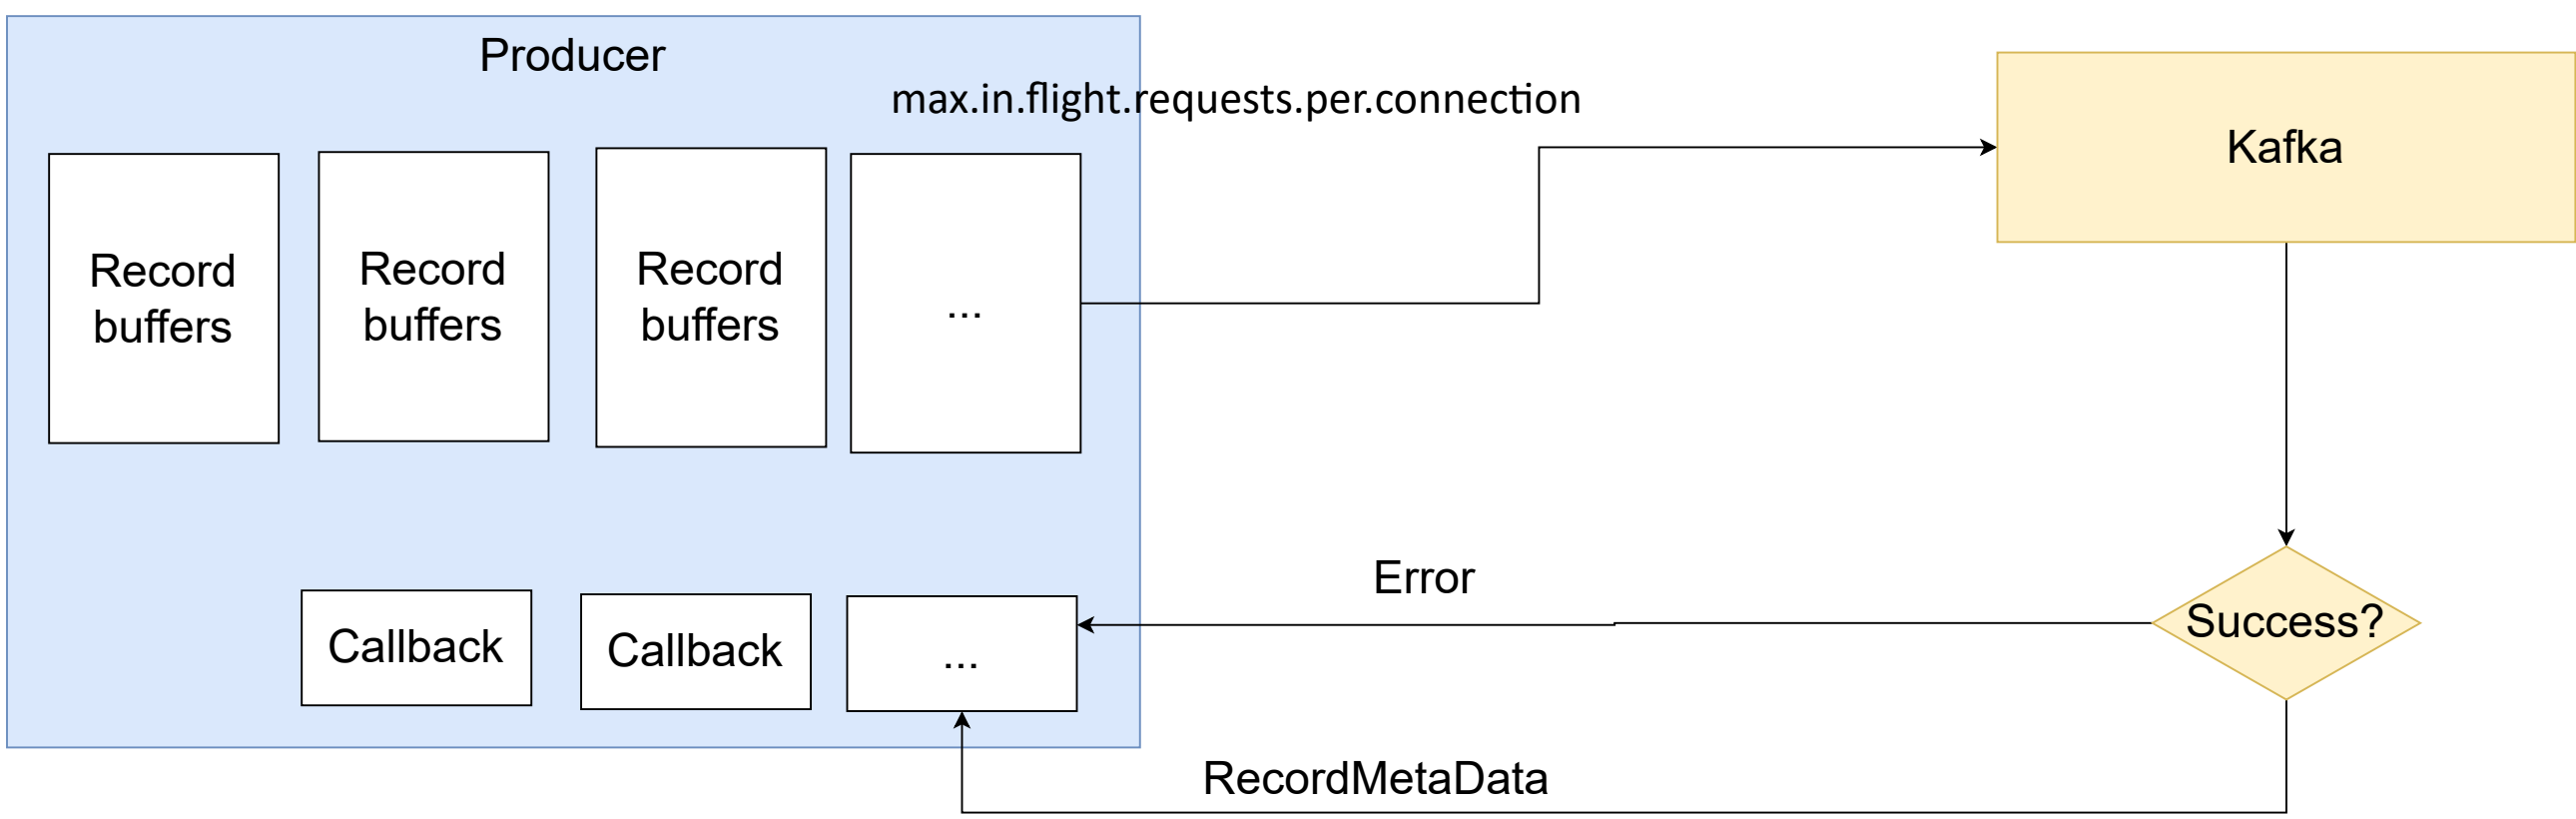
Producer Application Internals & Configurations



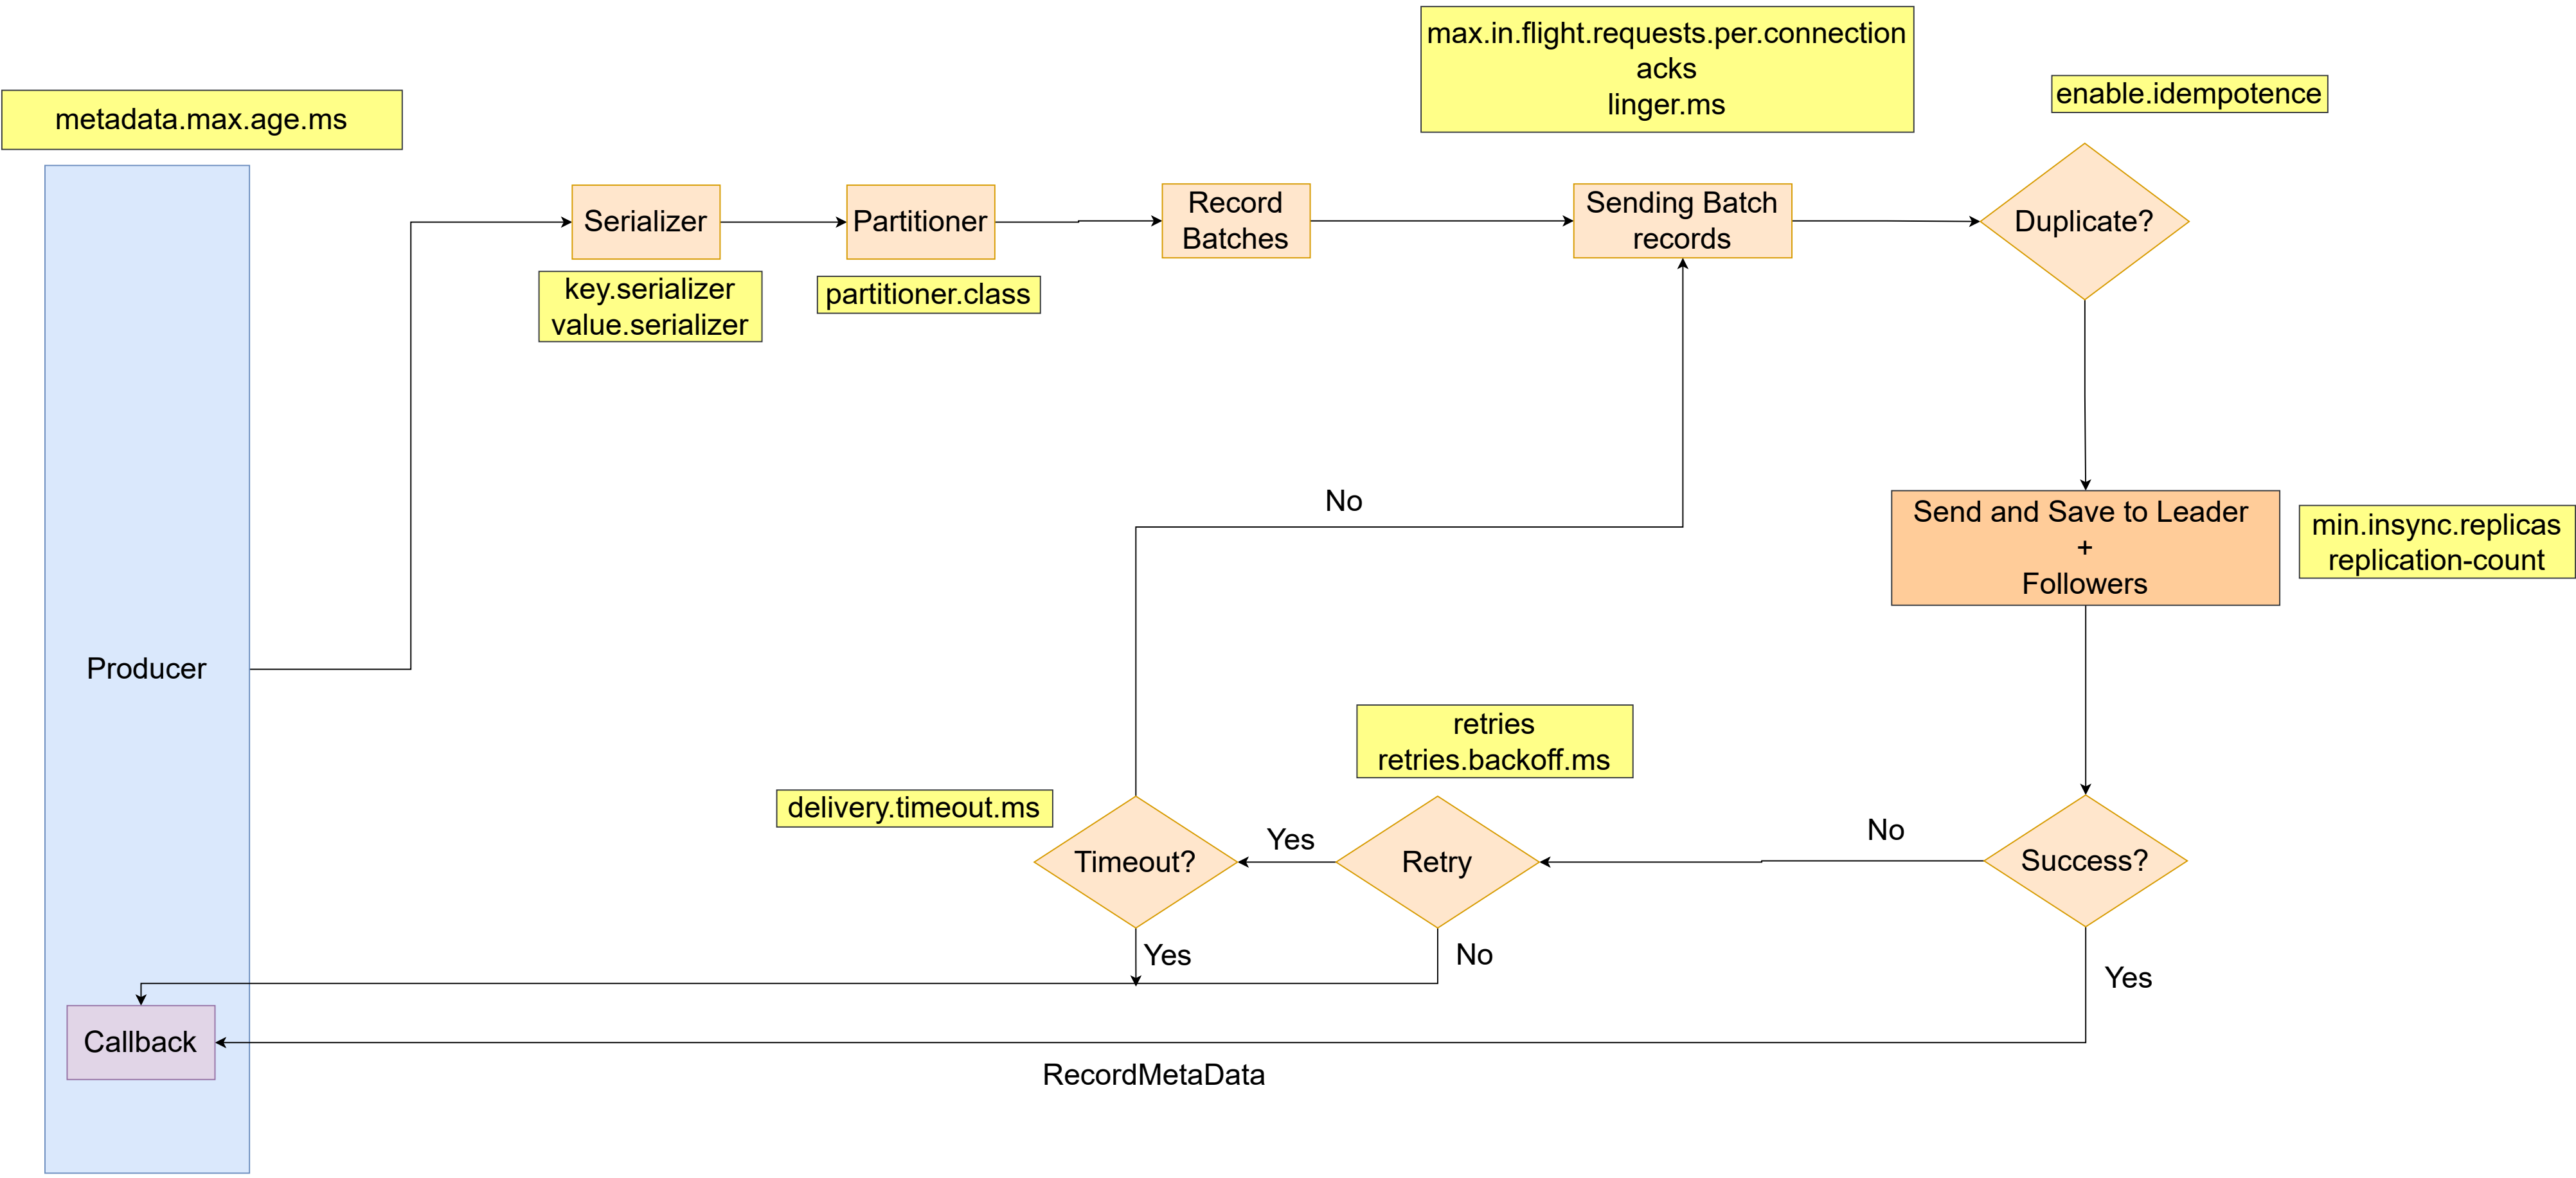
ack configuration



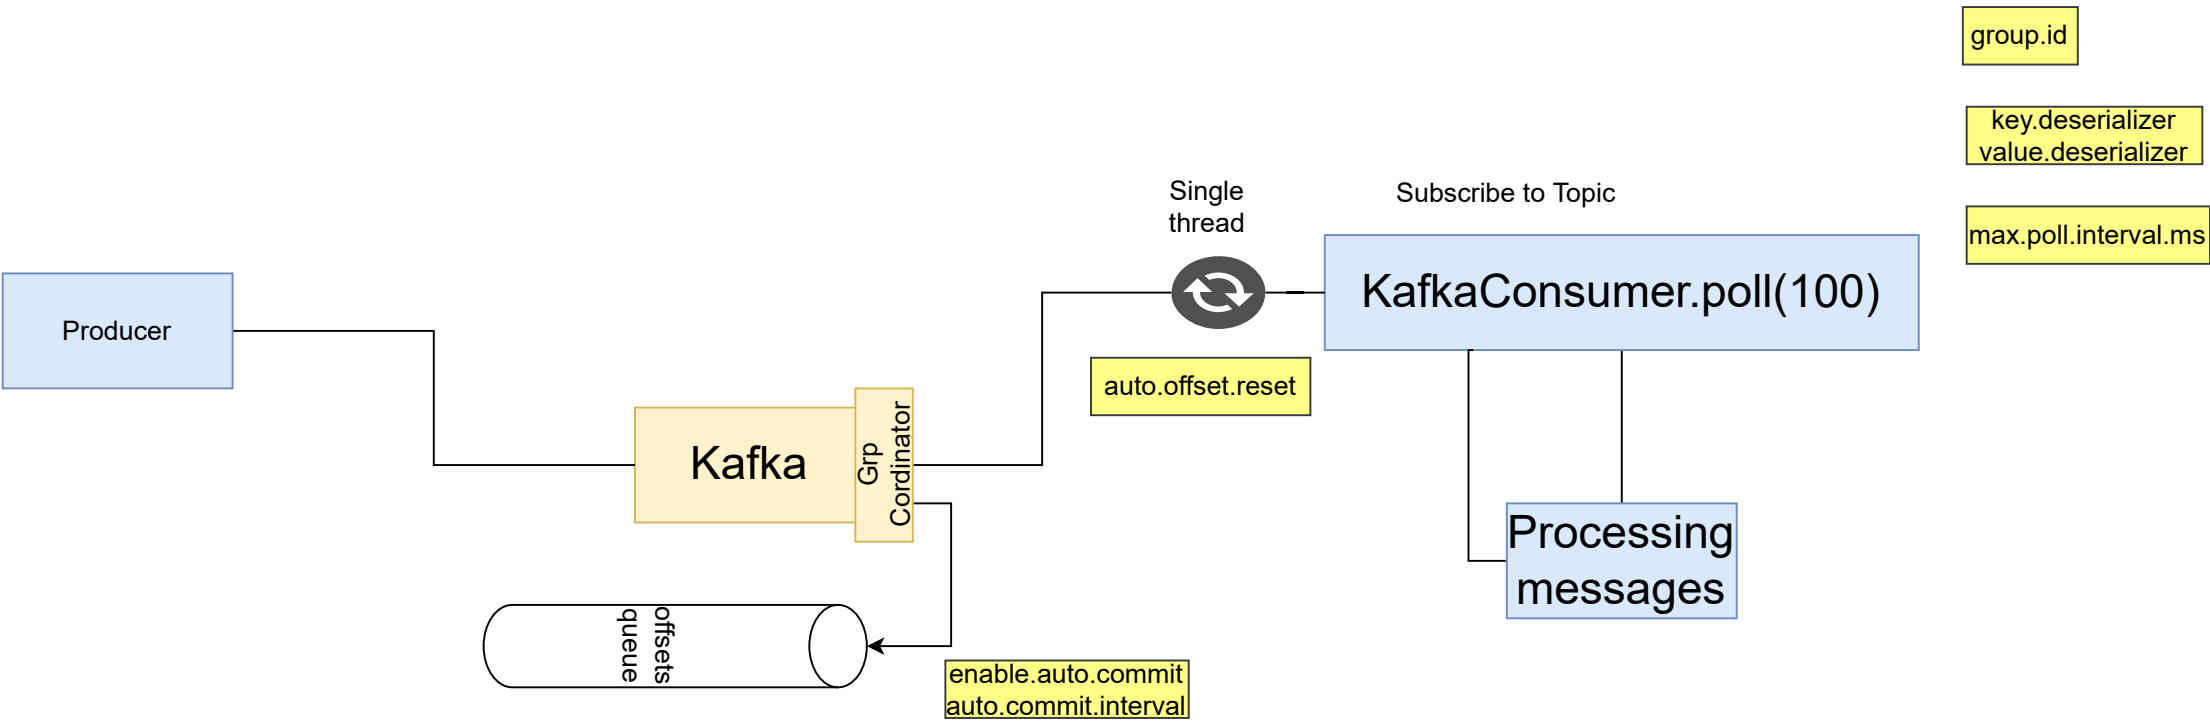
max.in.flight.requests.per.connection Configuration



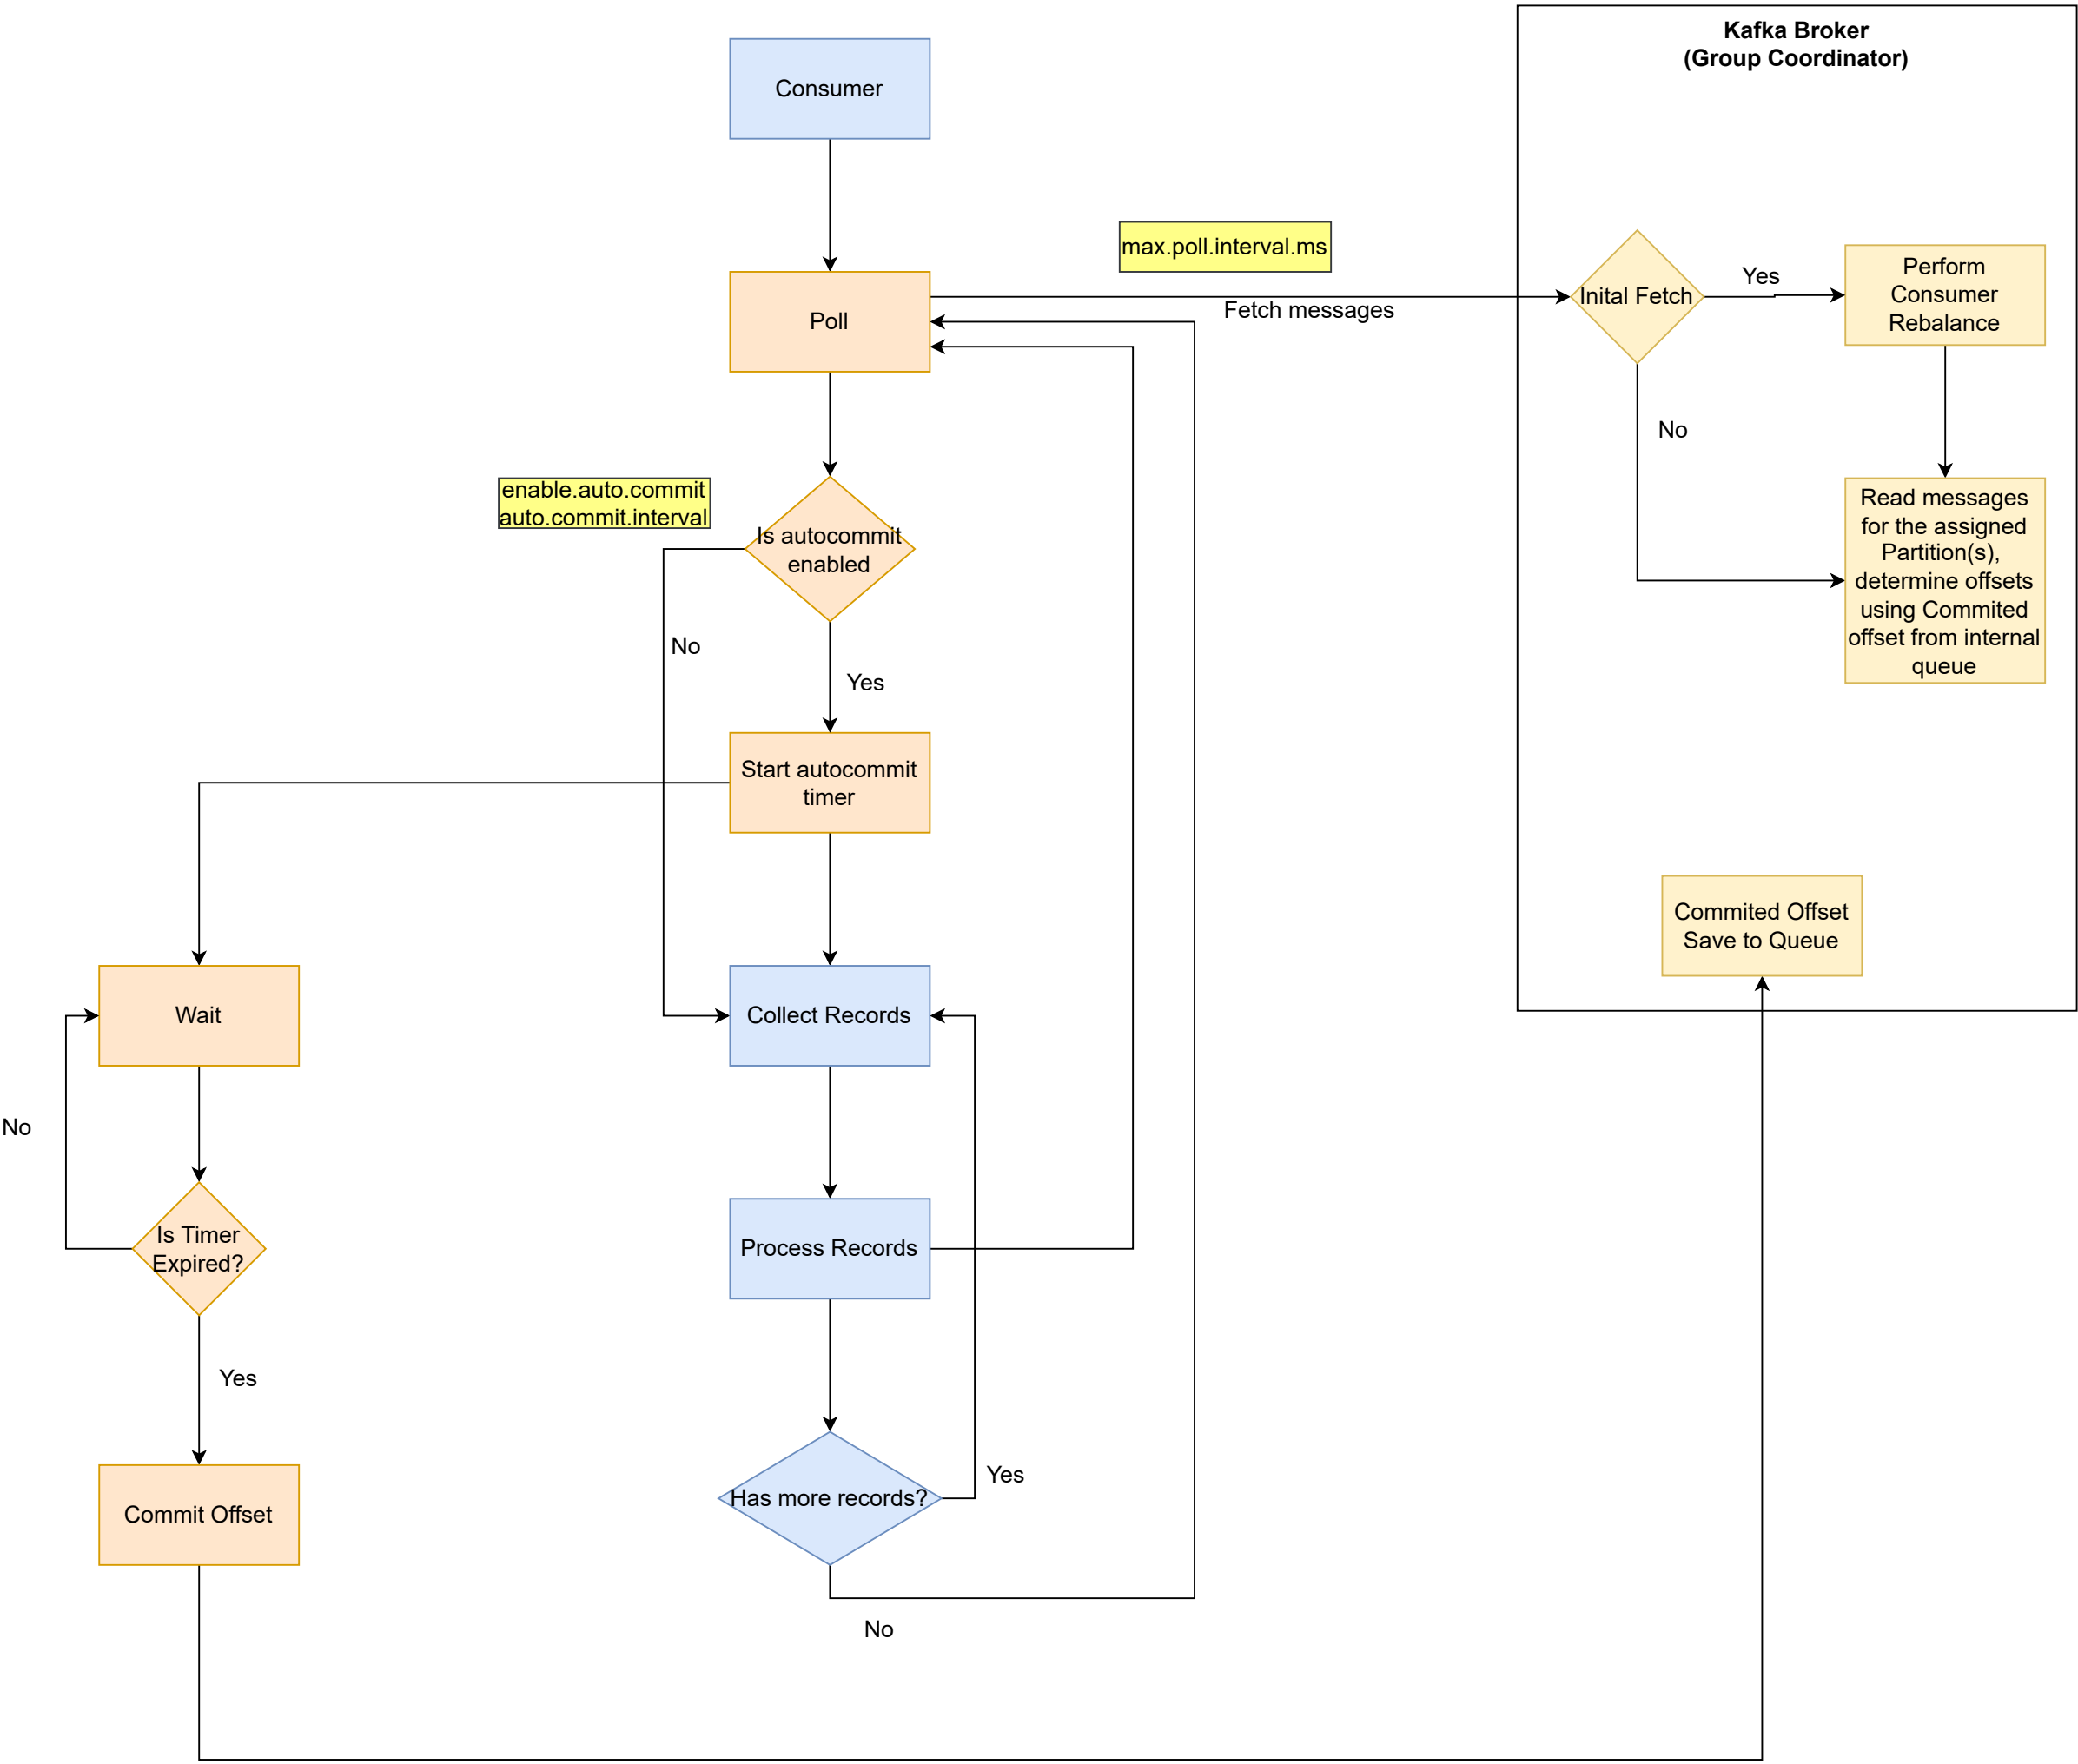
Consolidated Flow



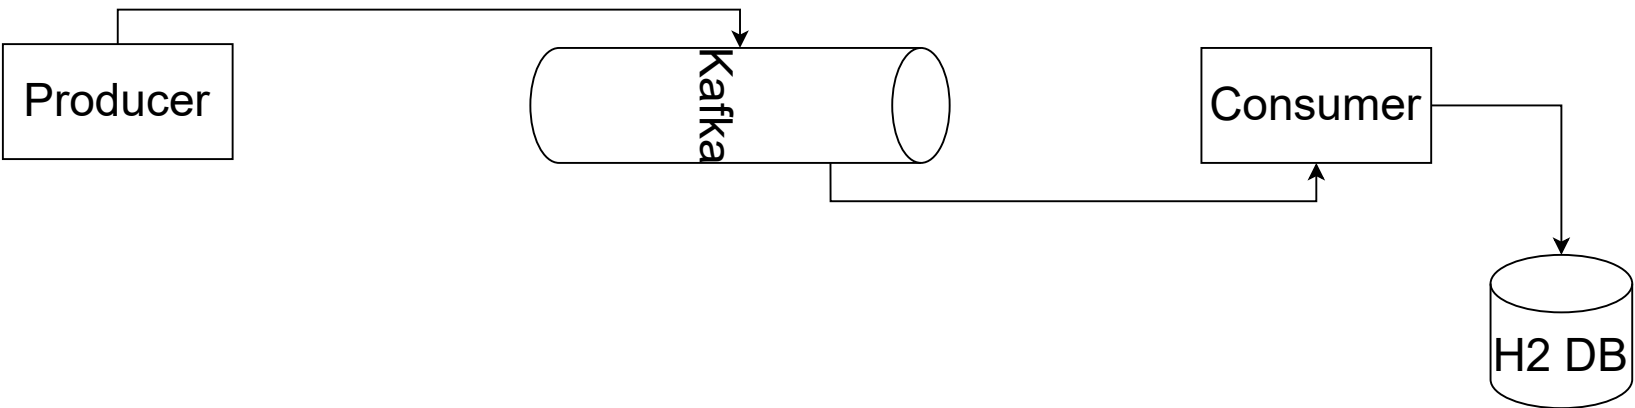
Consumer Pooling

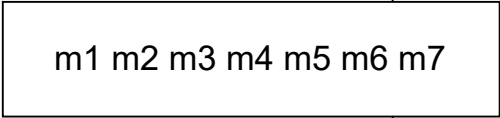
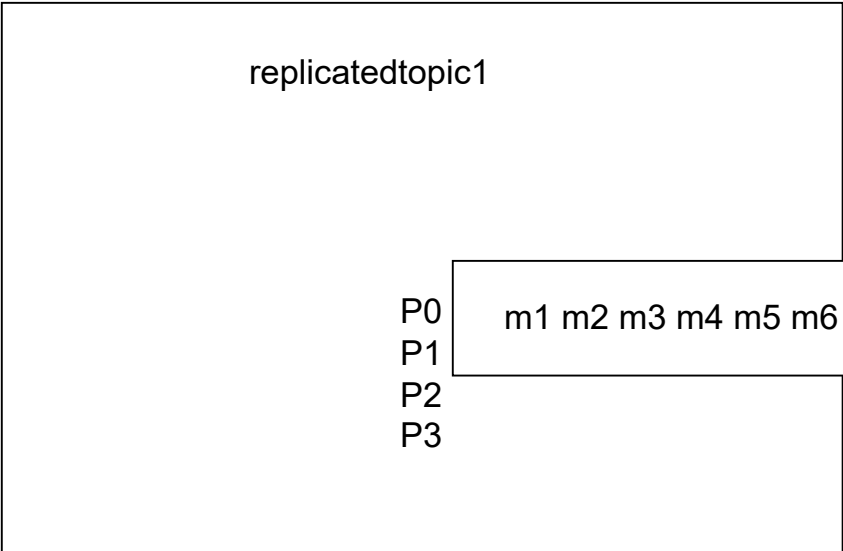
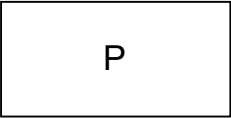


Consumer Internal Working



Persistency of read messages






P0
P1
P2
P3

m6 m7

m1 m2 m3 m4 m5


Current offset

C




Committed offset


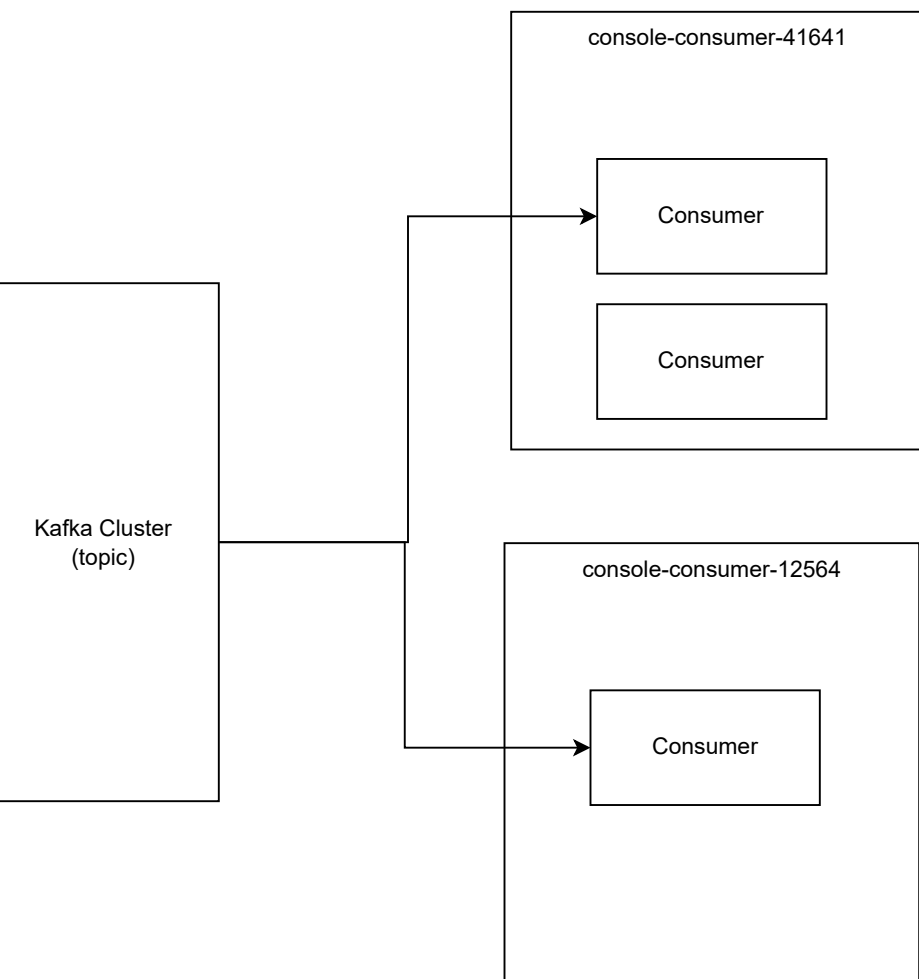
Partitions




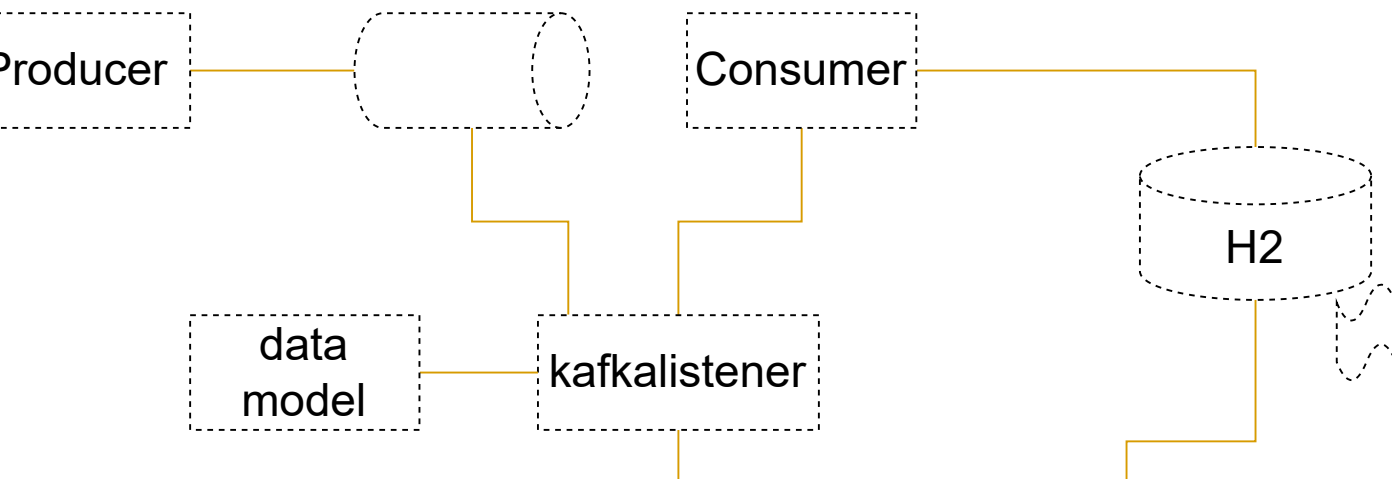
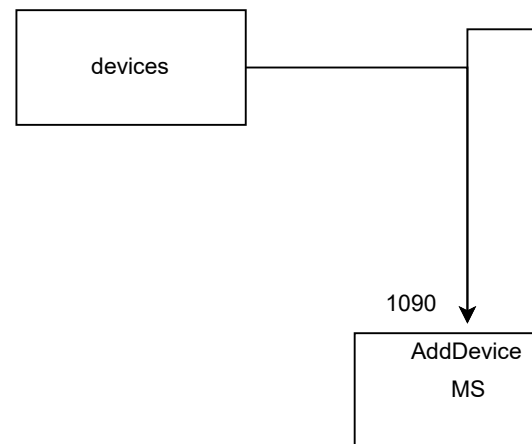
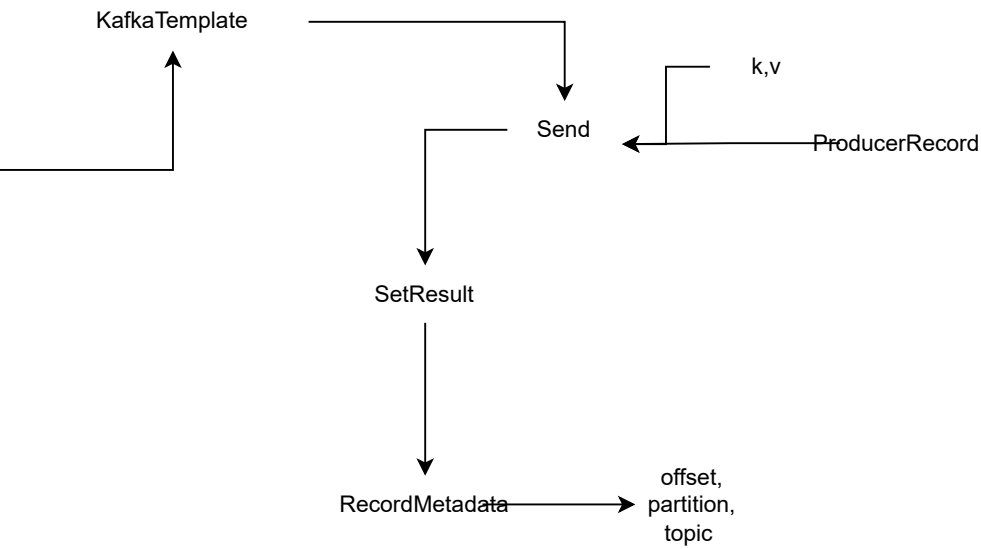
Config Kafkaconsu
Subscrib
Polling
Revocati
Assignme





umer
be
on
ent









me

message


Queue

Partion

Partion

Partion






message message

message


message message





thread 1



p1,
p2,
p3,
p4


Spring Kafka controllerd retry







Producer







Trust store



SSL / TLS certificate
public
key

public
key

server

client

root
CA
certs


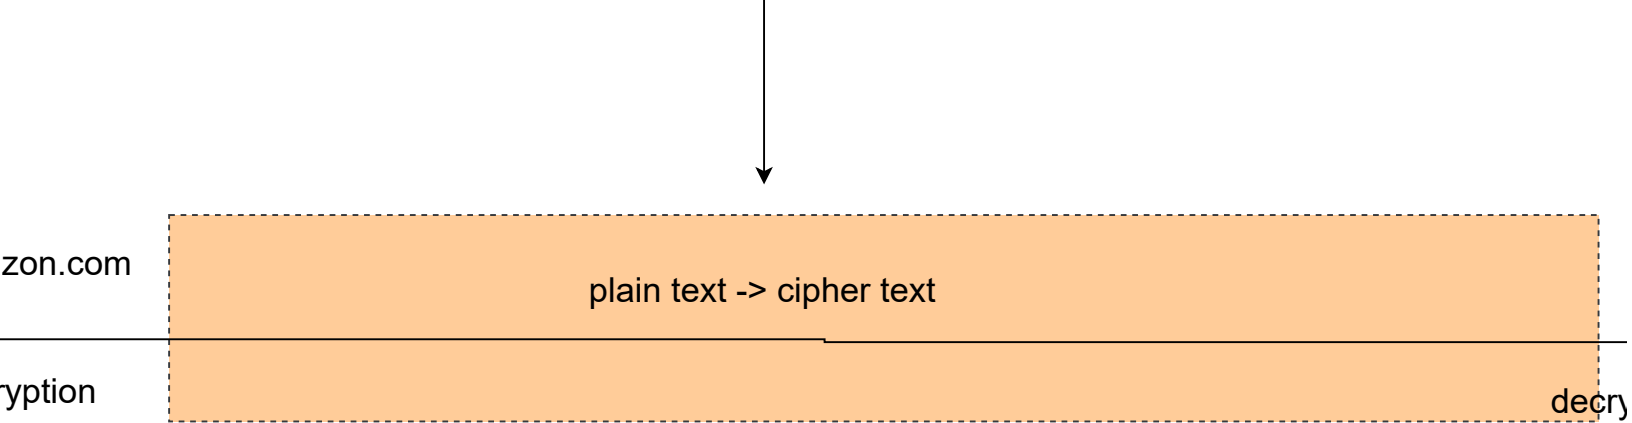
trusted store


signed certificate

keystores



Client





not
CA
certs


t CA cert

d certificate


CA

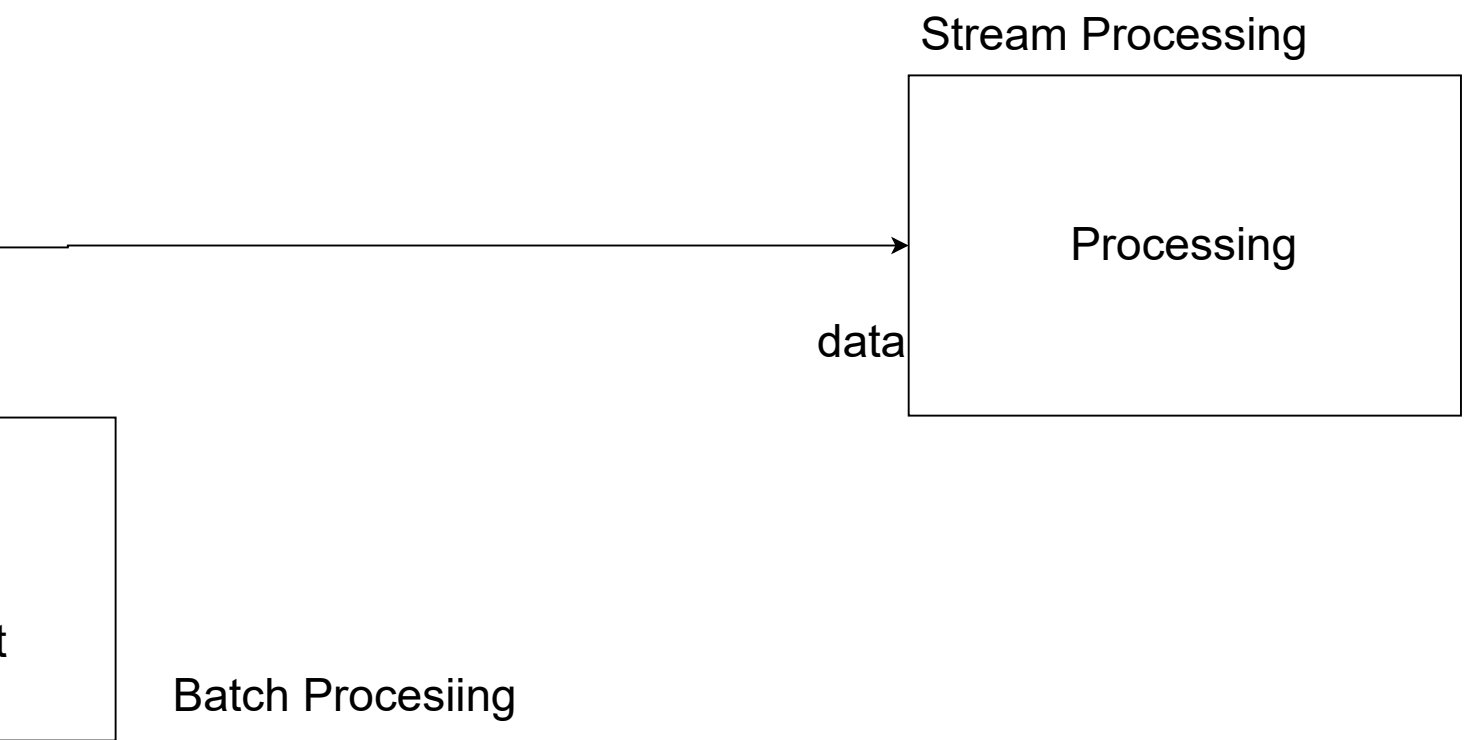
SSL certificates




SSL
- Encryption/ Decryption
- Authentication



StringBuilder



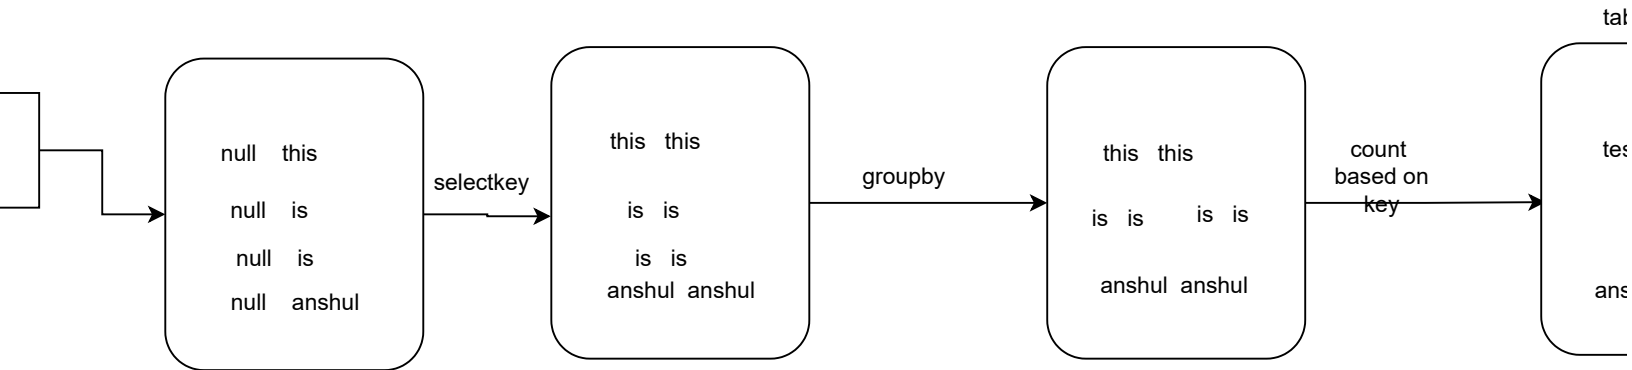




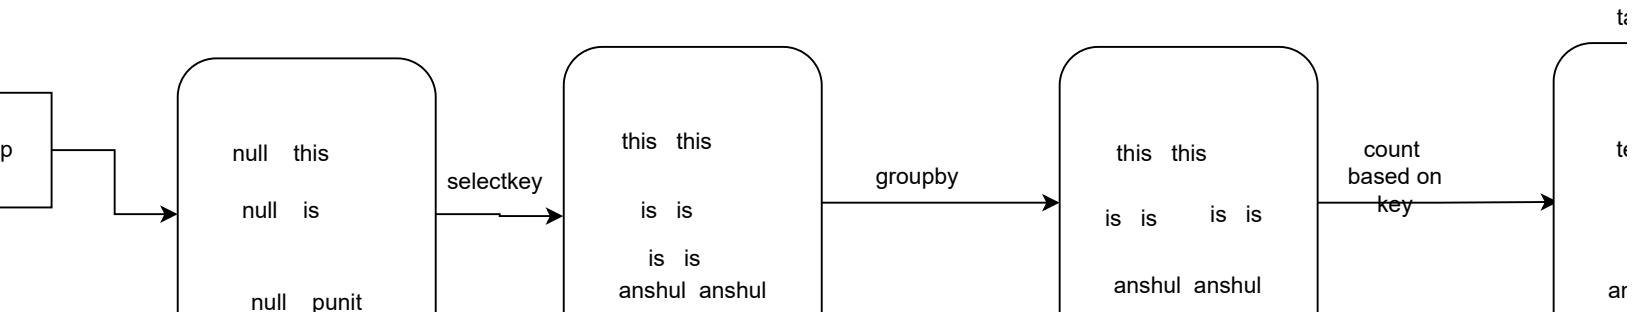
Processors (Source, Sink, processor)

Streaming
framework

hul



nit



ble

st 1

is 2


shul 1





able

est 1

is 2

hshul 1






record




flatMap




none



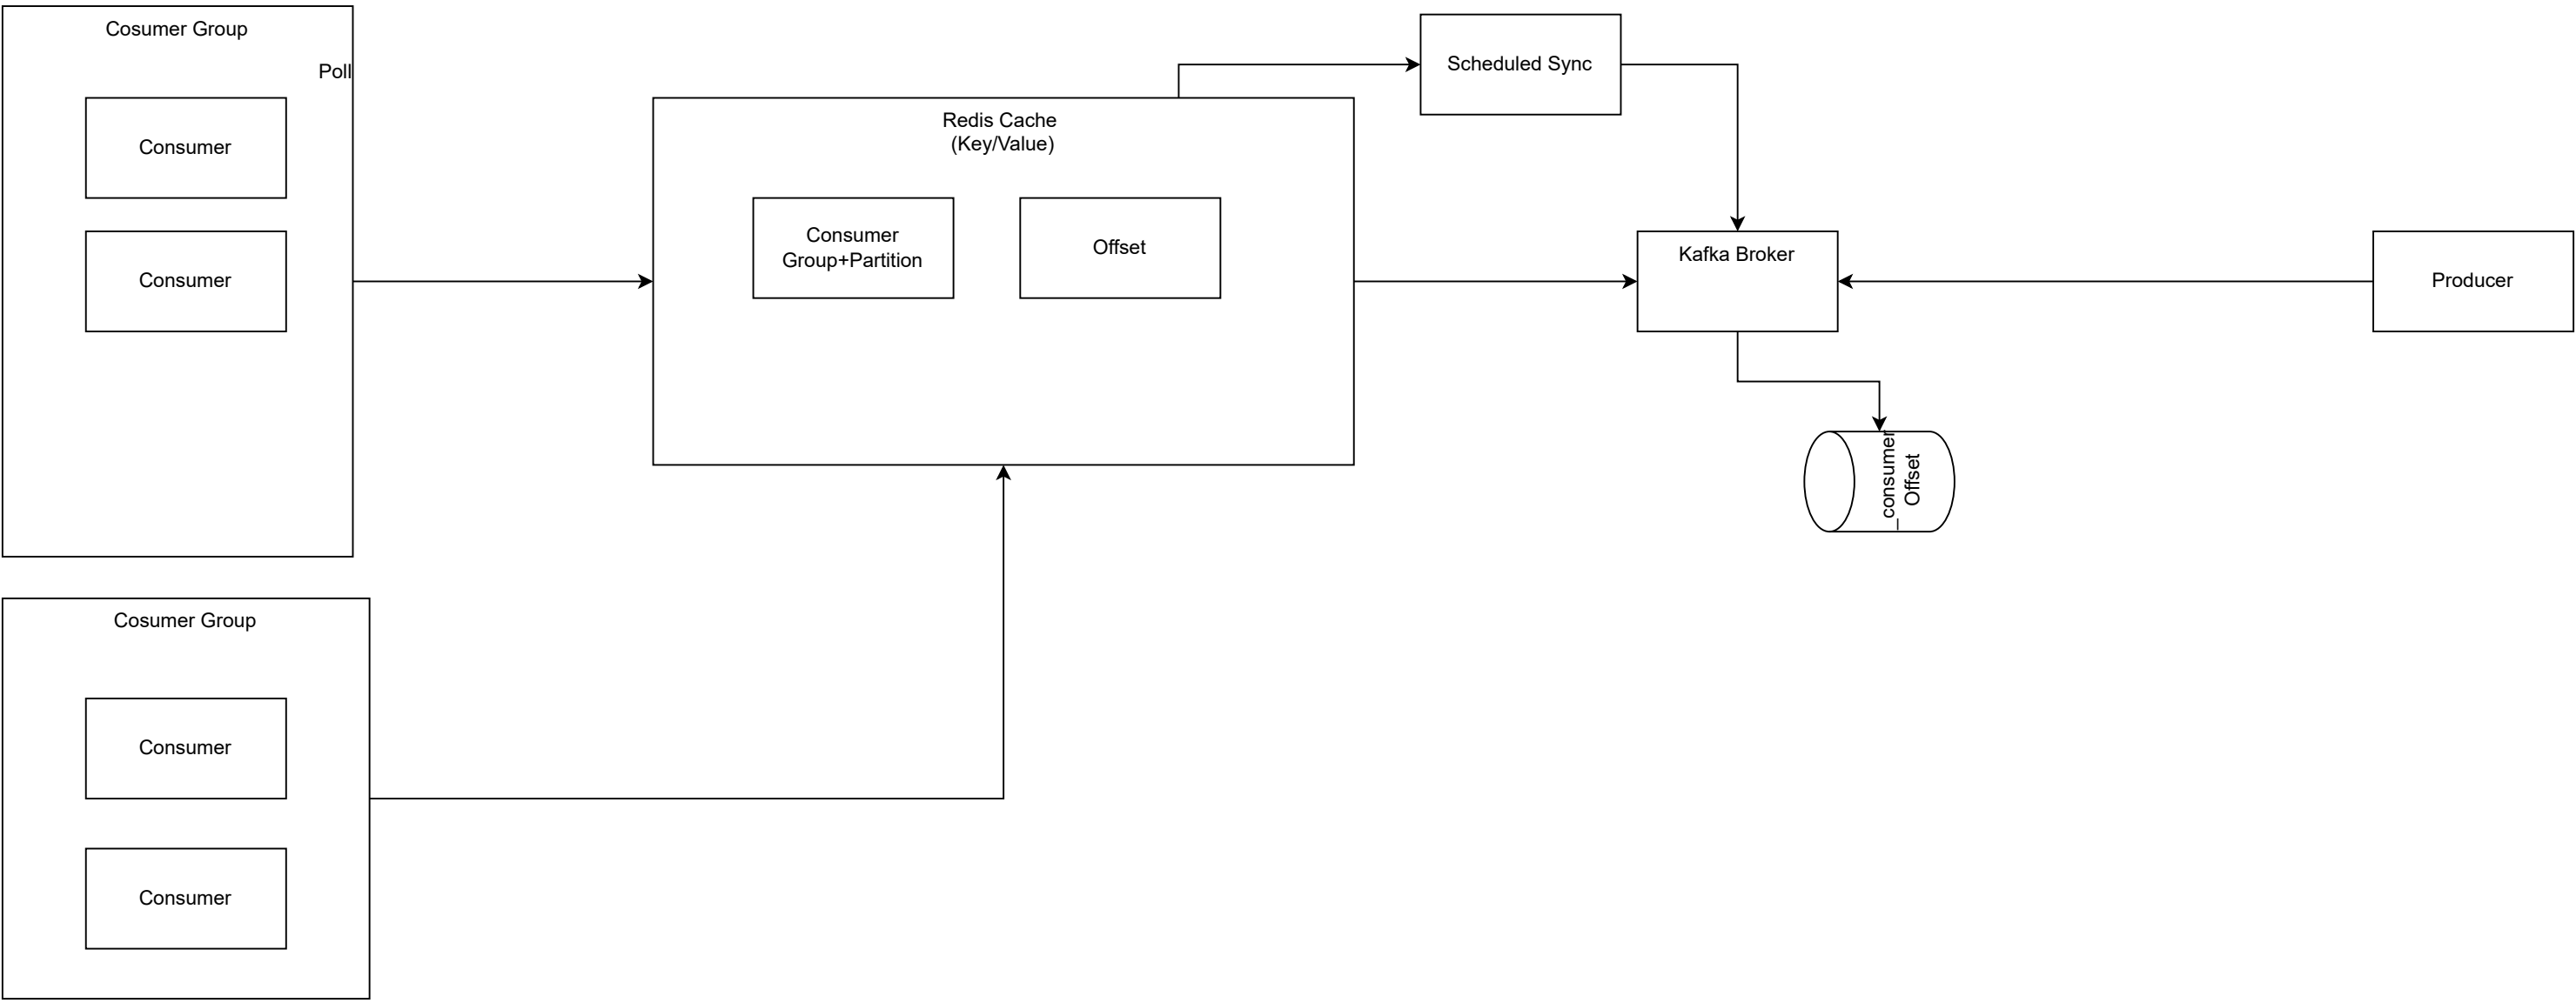
record




more than
one record



DIY Exercise - day 7

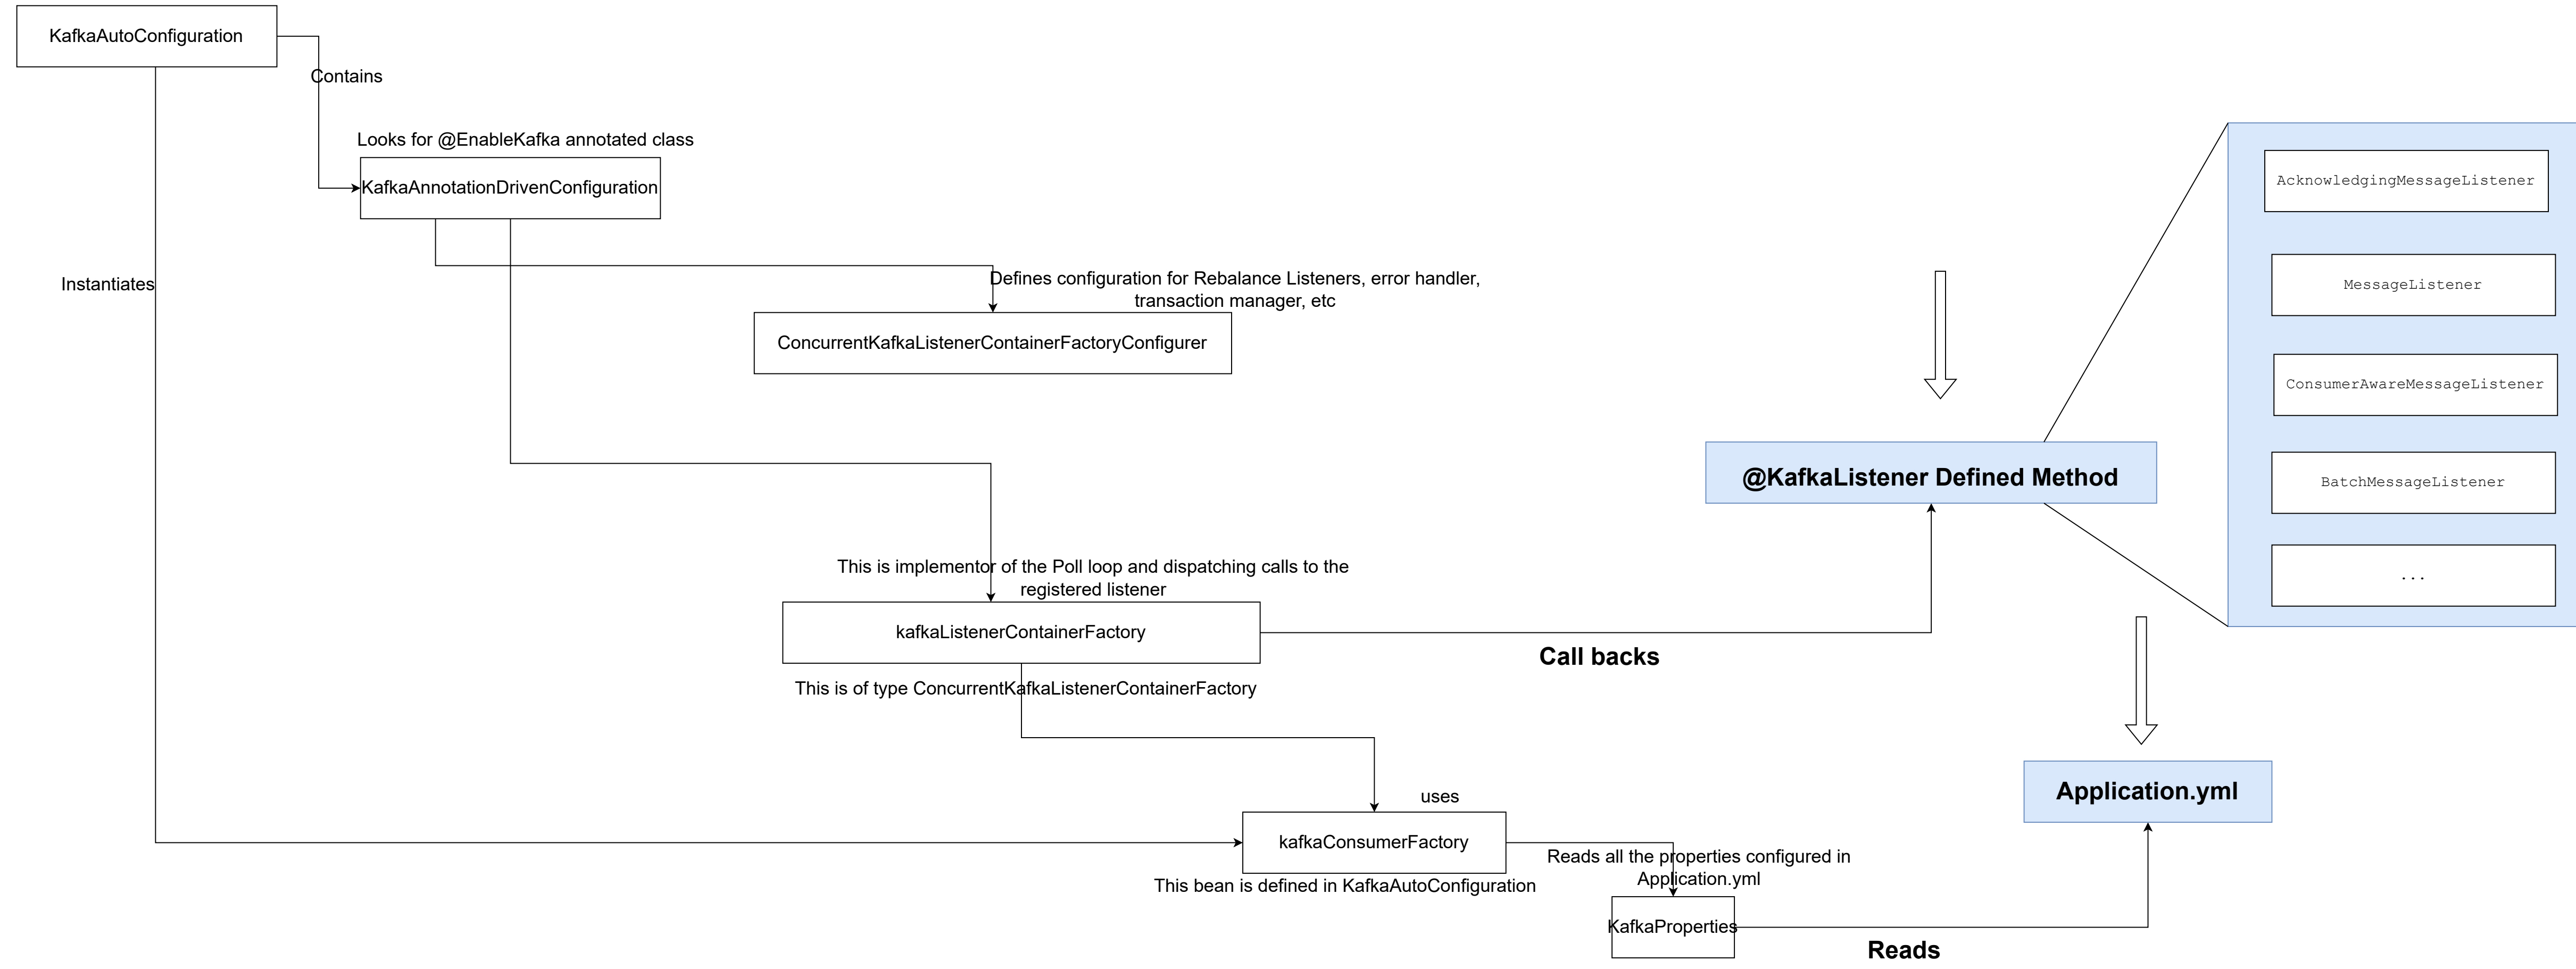


Producer Retry Mechanism




Spring Kafka Consumer Internals

Instantiate needed beans and initiates Other objects

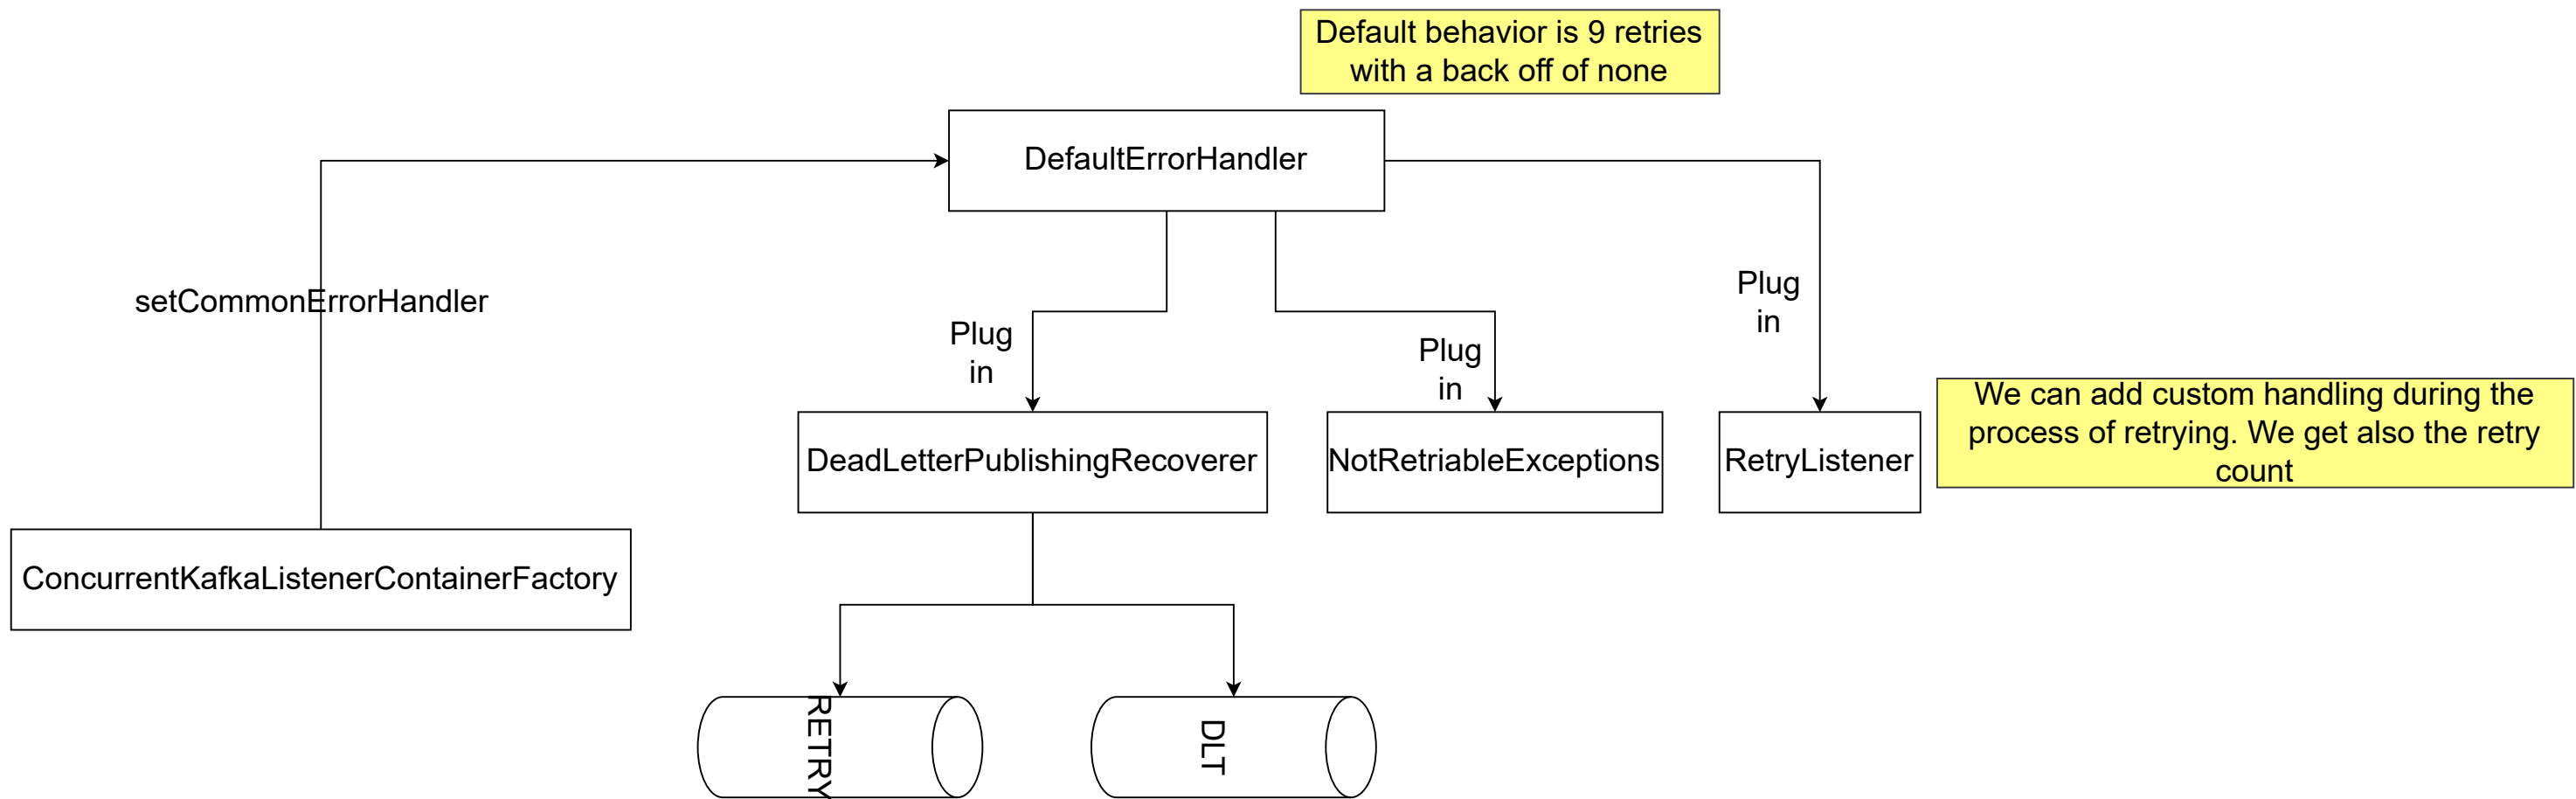


Producer Test

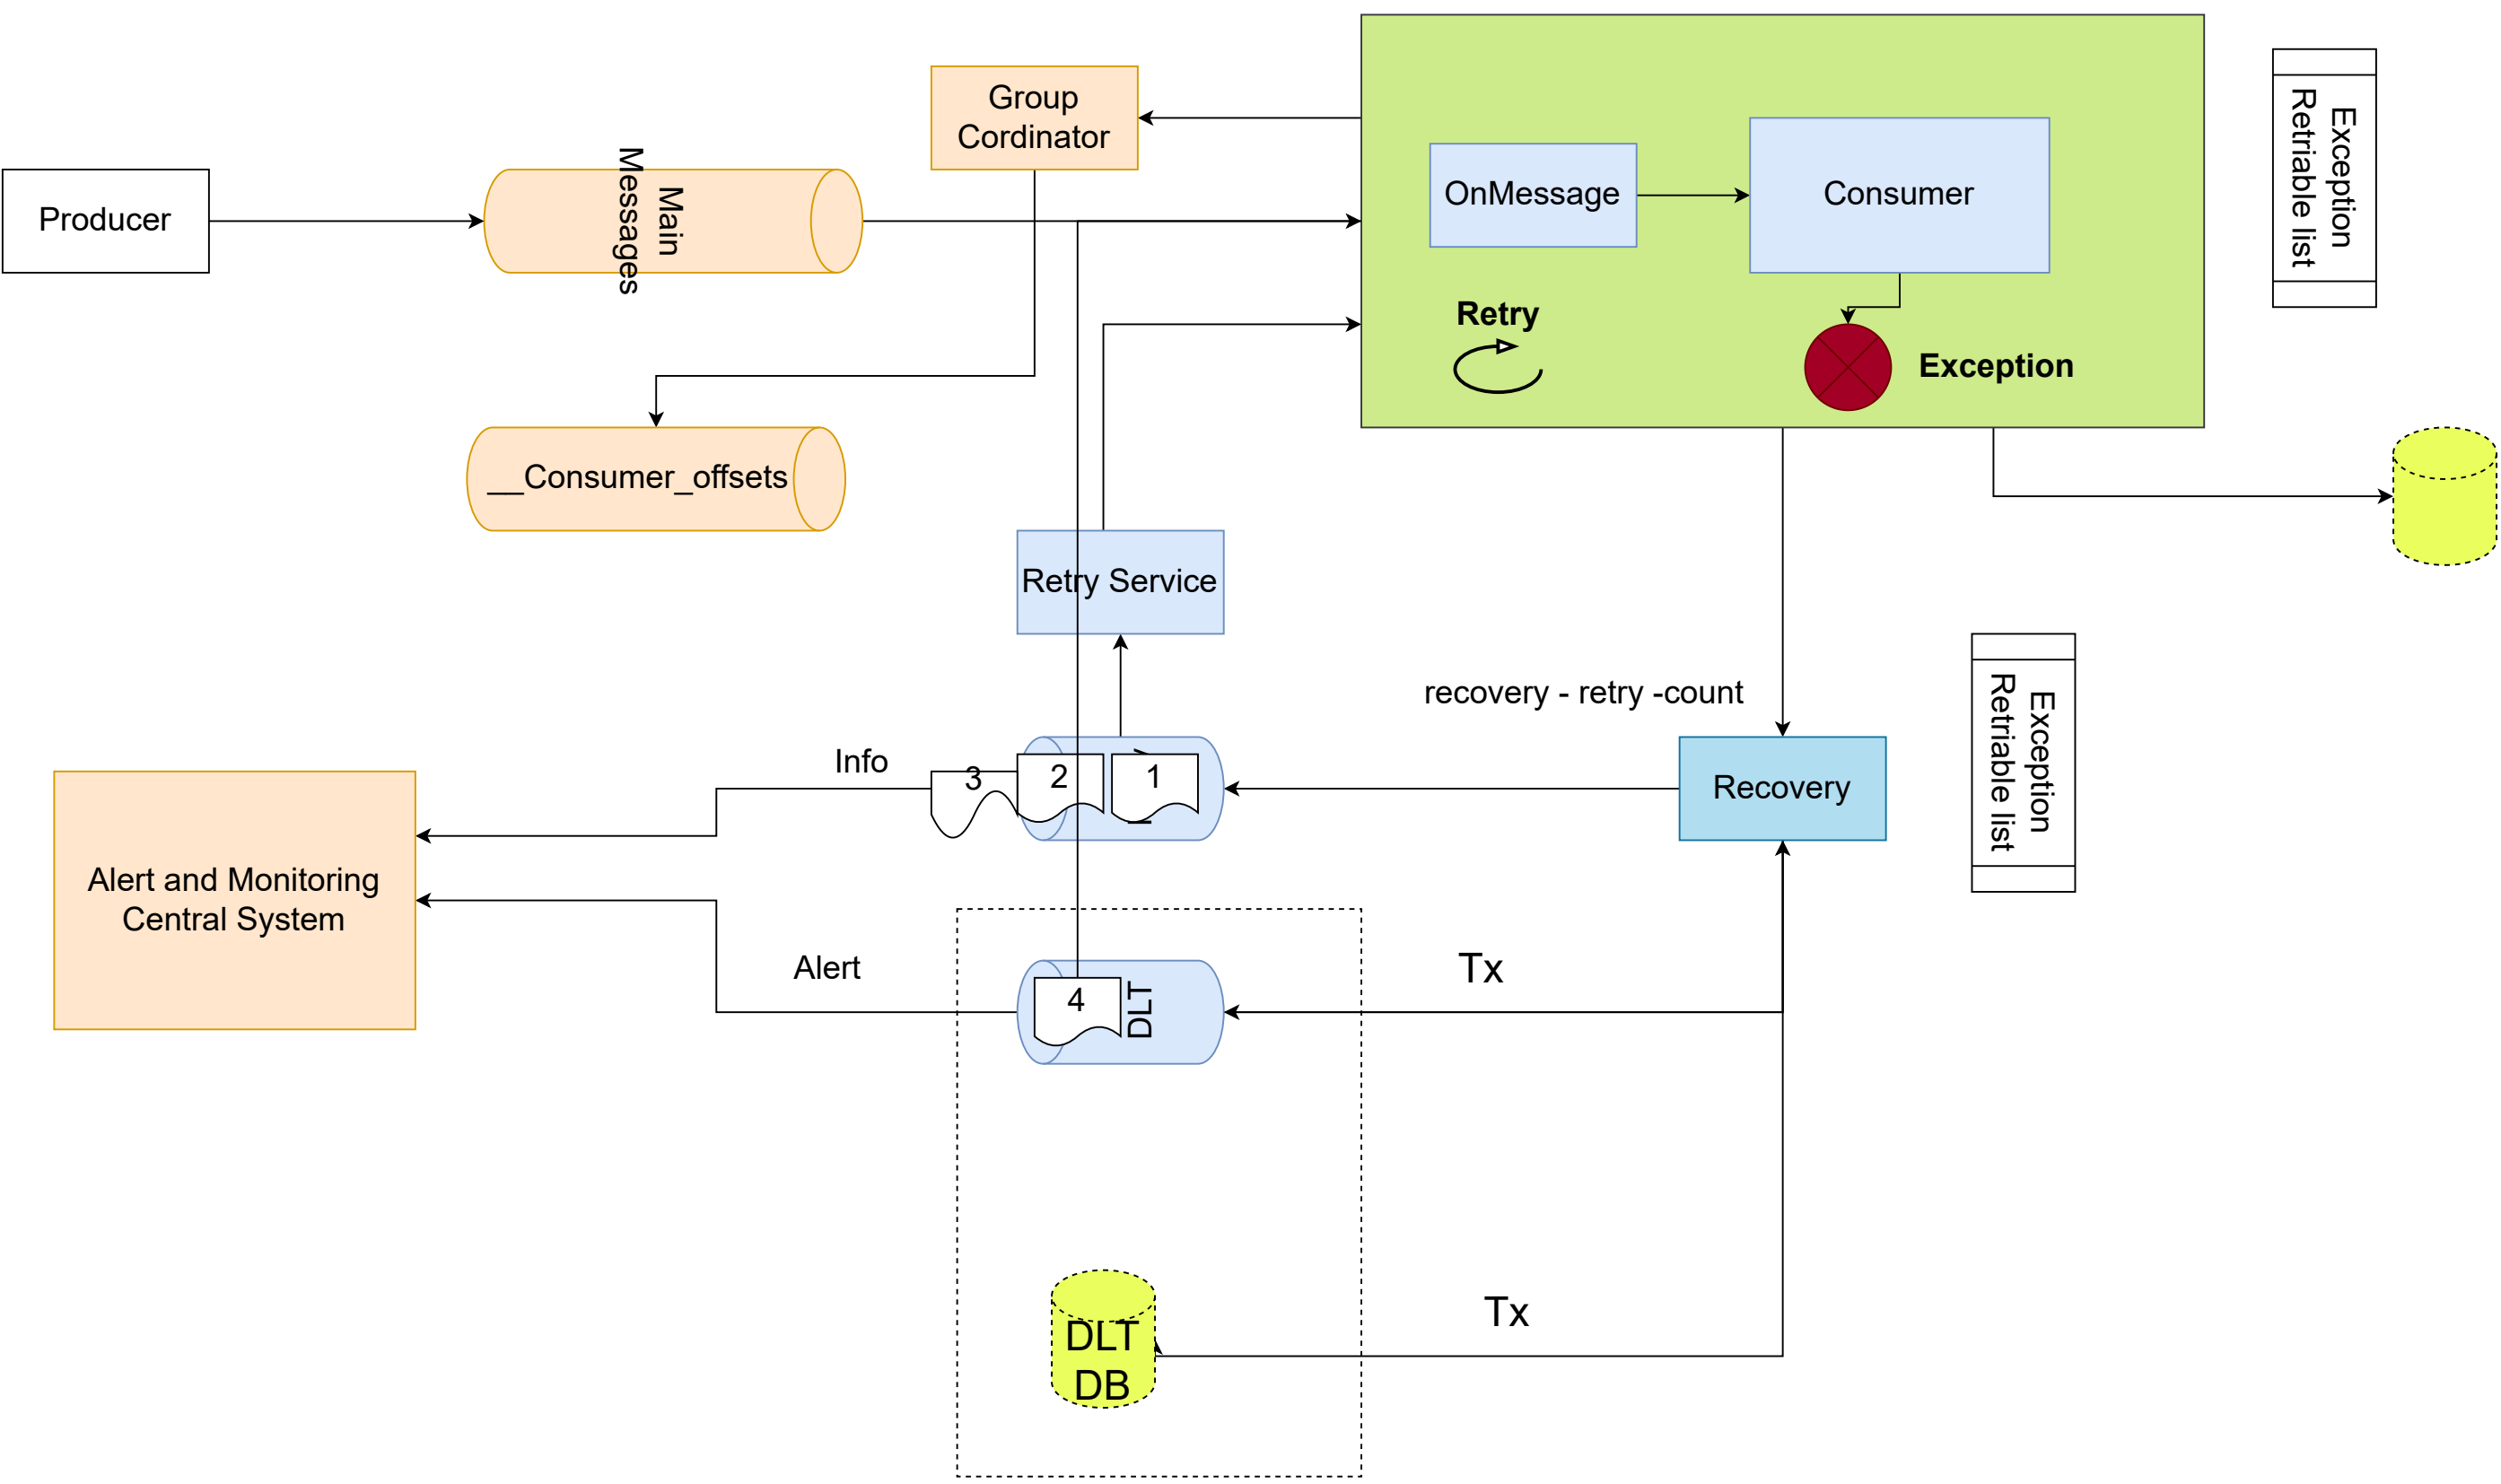


Consumer Error Handling & Recovery Mechanism


Spring Kafka Consumer Classes




Recovery



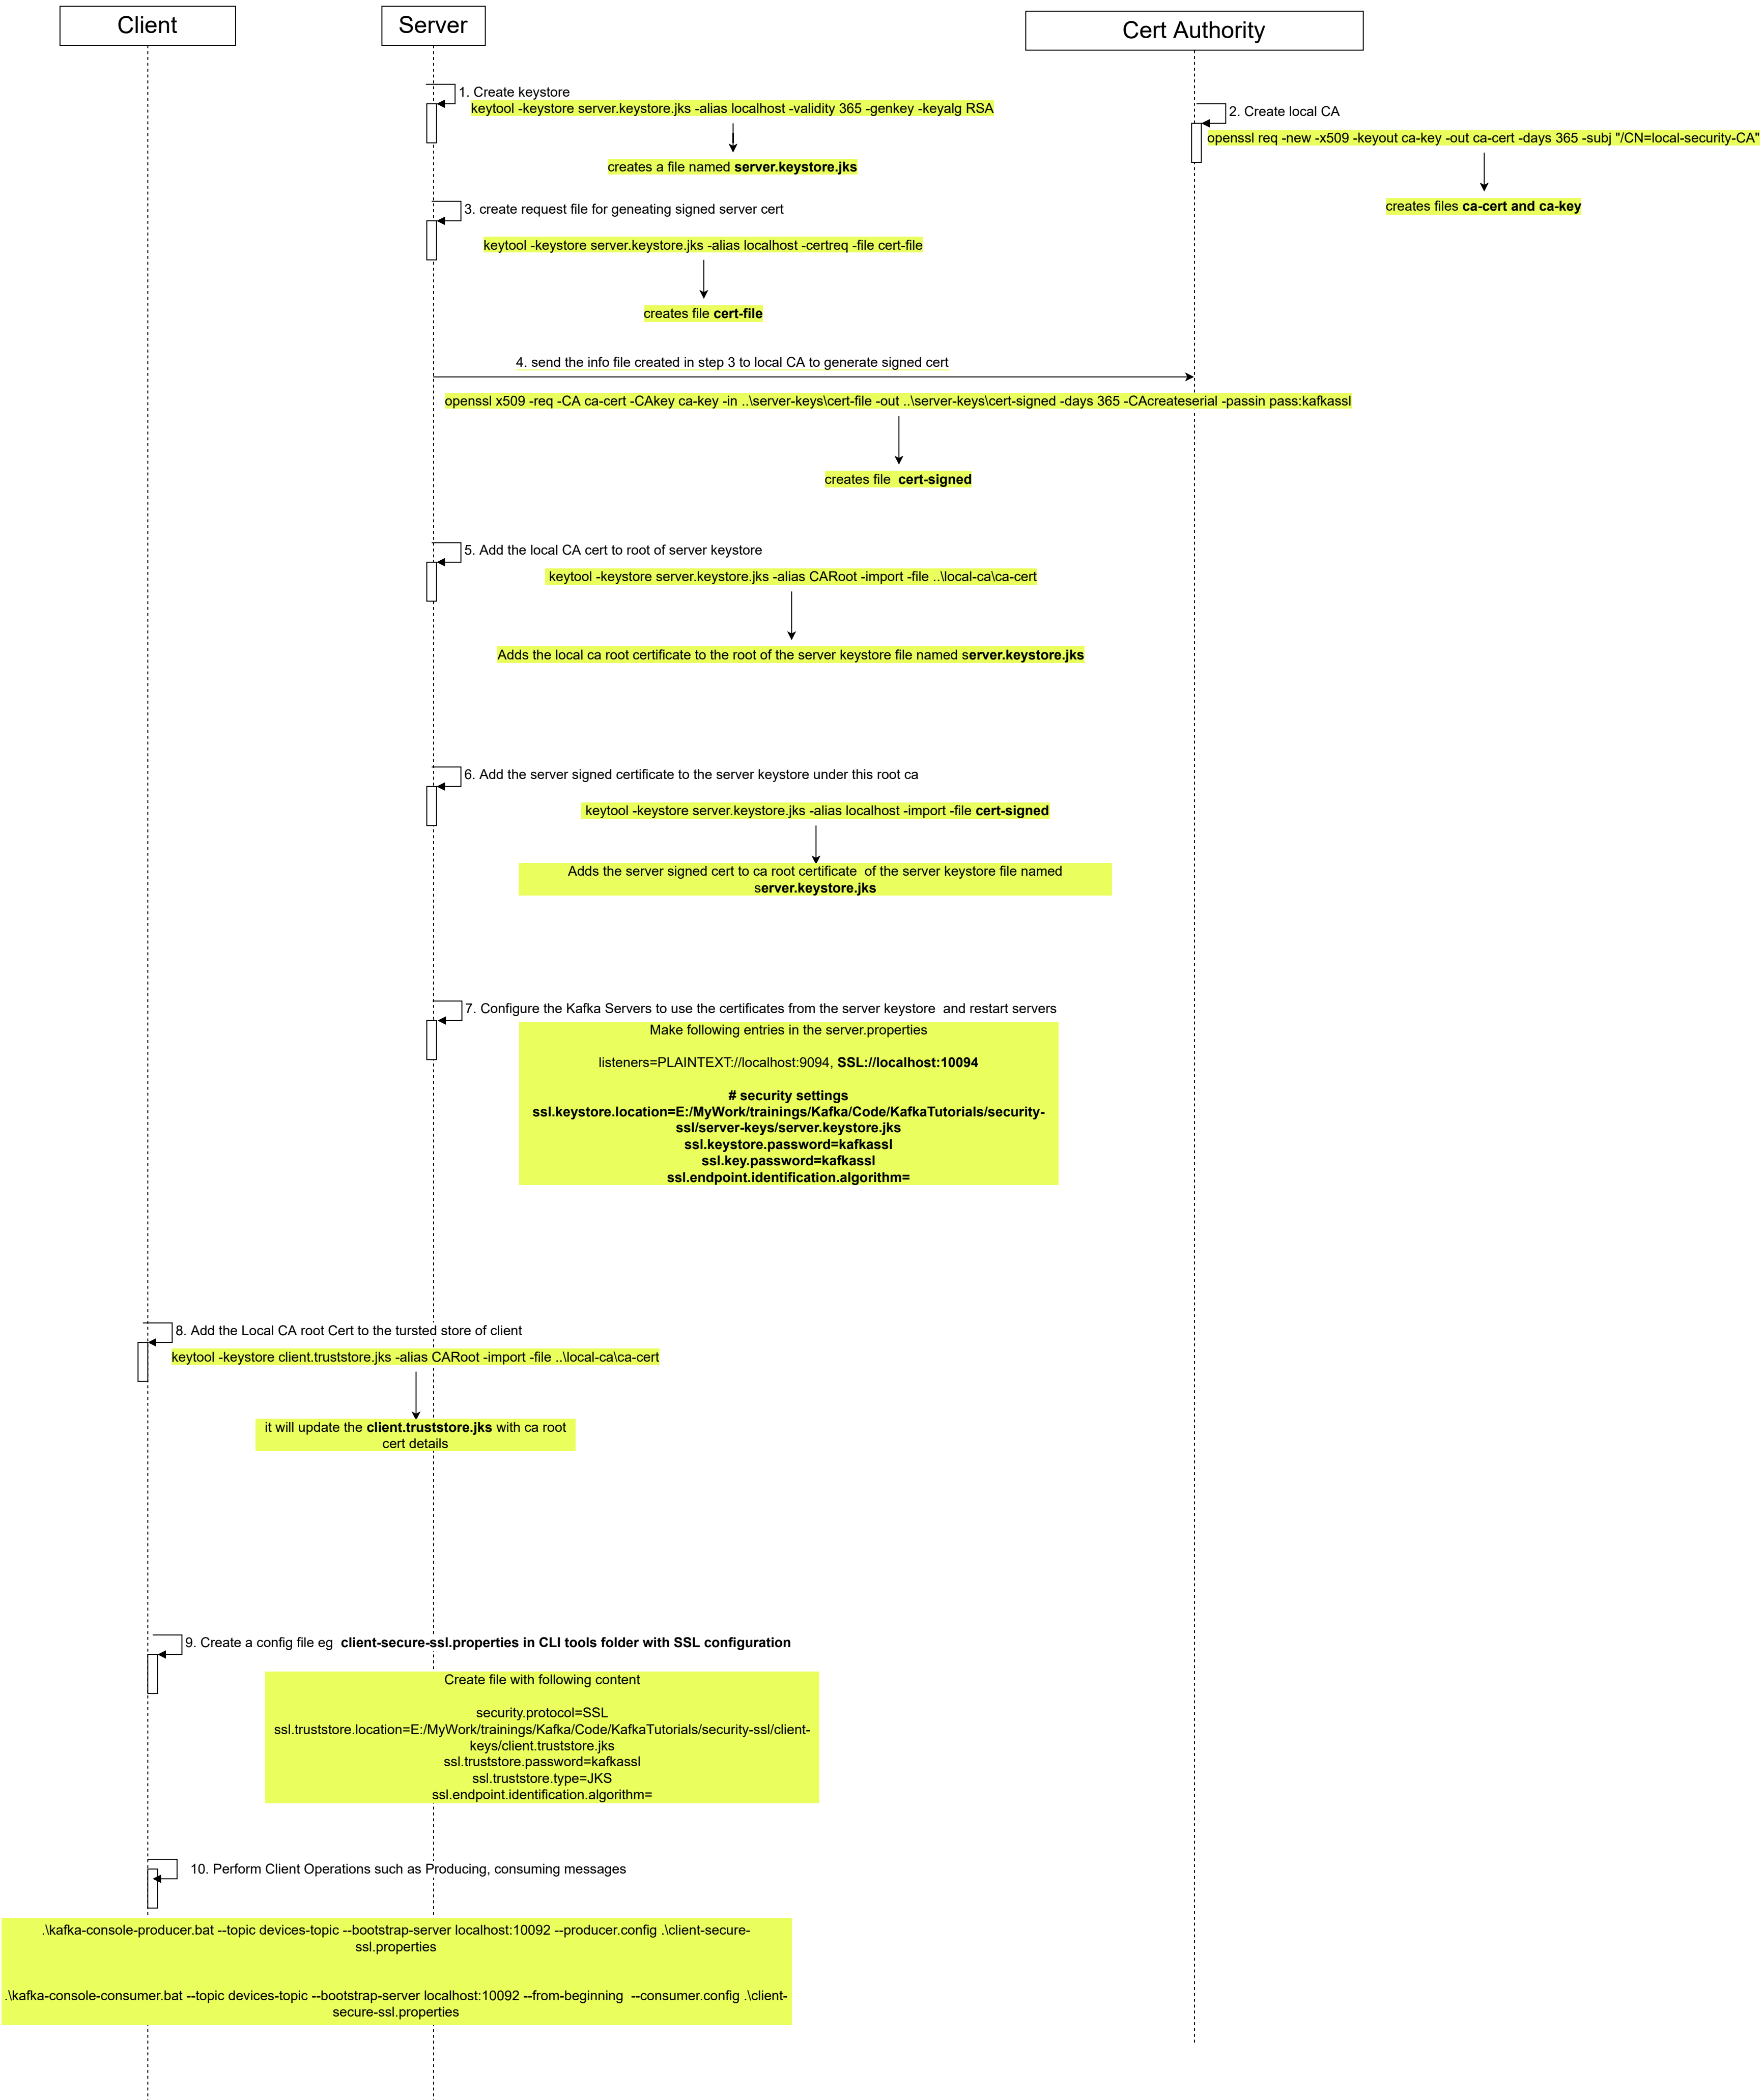
Communications in Kafka Infra To Be Secured




Securing Kafka Client and Server using SSL



SSL Setup for securing Kafka Server and Clients - 1 Way




2101 100



User creation

ACL creation

SASL Topology



Important commands

Admin	<code>./bin/kafka-configs.sh --zookeeper localhost:2182 --zk-tls-config-file ./bin/zookeeper-client.properties --alter --add-conf</code>
Producer	<code>./bin/kafka-configs.sh --zookeeper localhost:2182 --zk-tls-config-file ./bin/zookeeper-client.properties --alter --add-conf</code>
Consumer	<code>./bin/kafka-configs.sh --zookeeper localhost:2182 --zk-tls-config-file ./bin/zookeeper-client.properties --alter --add-conf</code>
Producer	<code>./bin/kafka-acls.sh --authorizer-properties zookeeper.connect=localhost:2182 --zk-tls-config-file ./bin/zookeeper-client.properties --add-acl</code>
Consumer	<code>./bin/kafka-acls.sh --authorizer-properties zookeeper.connect=localhost:2182 --zk-tls-config-file ./bin/zookeeper-client.properties --add-acl</code>
Consumer to group	<code>./bin/kafka-acls.sh --authorizer-properties zookeeper.connect=localhost:2182 --zk-tls-config-file ./bin/zookeeper-client.properties --add-acl</code>

fig 'SCRAM-SHA-512=[password=kafkassl]' --entity-type users --entity-name broker-admin

fig 'SCRAM-SHA-512=[password=kafkassl]' --entity-type users --entity-name kafkaprod

fig 'SCRAM-SHA-512=[password=kafkassl]' --entity-type users --entity-name kafkacon

ookeeper-client.properties --add --allow-principal User:kafkaprod --topic secure-topic --operation WRITE --operation DESCRIBE --t

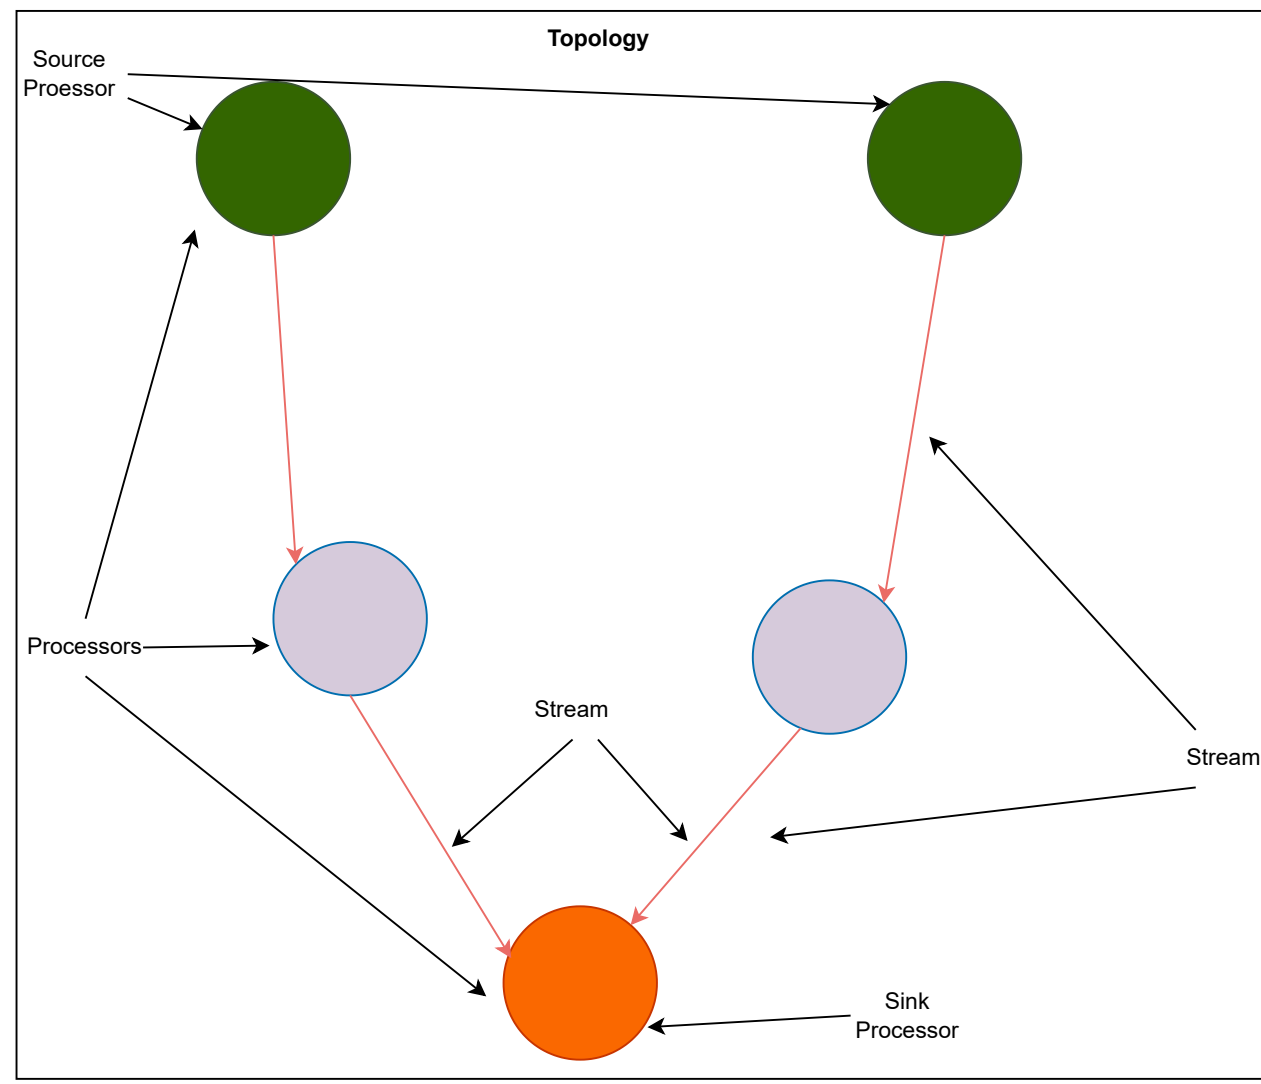

ookeeper-client.properties --add --allow-principal User:kafkacon --topic secure-topic --operation READ --operation DESCRIBE --oper:

ookeeper-client.properties --add --allow-principal User:kafkacon --group secure-group --operation READ

operation DESCRIBECONFIGS


ation DESCRIBECONFIGS

Kafka Streams




- 1. create StreamBuilder
- 2. Source Stream processor
- 3. Processing processors
- 4. build the builder to create topology
- 5. KafkaStreams create object (topology and config)
- 6. use the method of kafkaStreams to execute the topology


Basic Topology



WordCount Topology




KStream




Create Stream from Kafka Topic-
StreamBuilder.stream

KStream




Change all values to lower case for predictable calculation-
mapValues

KStream




Split each word in value as seperate record-
flatMapValues

KStream




Update the null keys with same as value-
selectKey

KGroupedStream




Group the records based on key-
groupByKey

KTable



Count the records in each group-
count




Write into new topic key value pairs having word and its count
to

user, color




Stream the user and color from kafka

user, color




Filter out colors which are only RED, GREEN or
BLUE only

color, user




Change key to be color

color, user



groupbykey


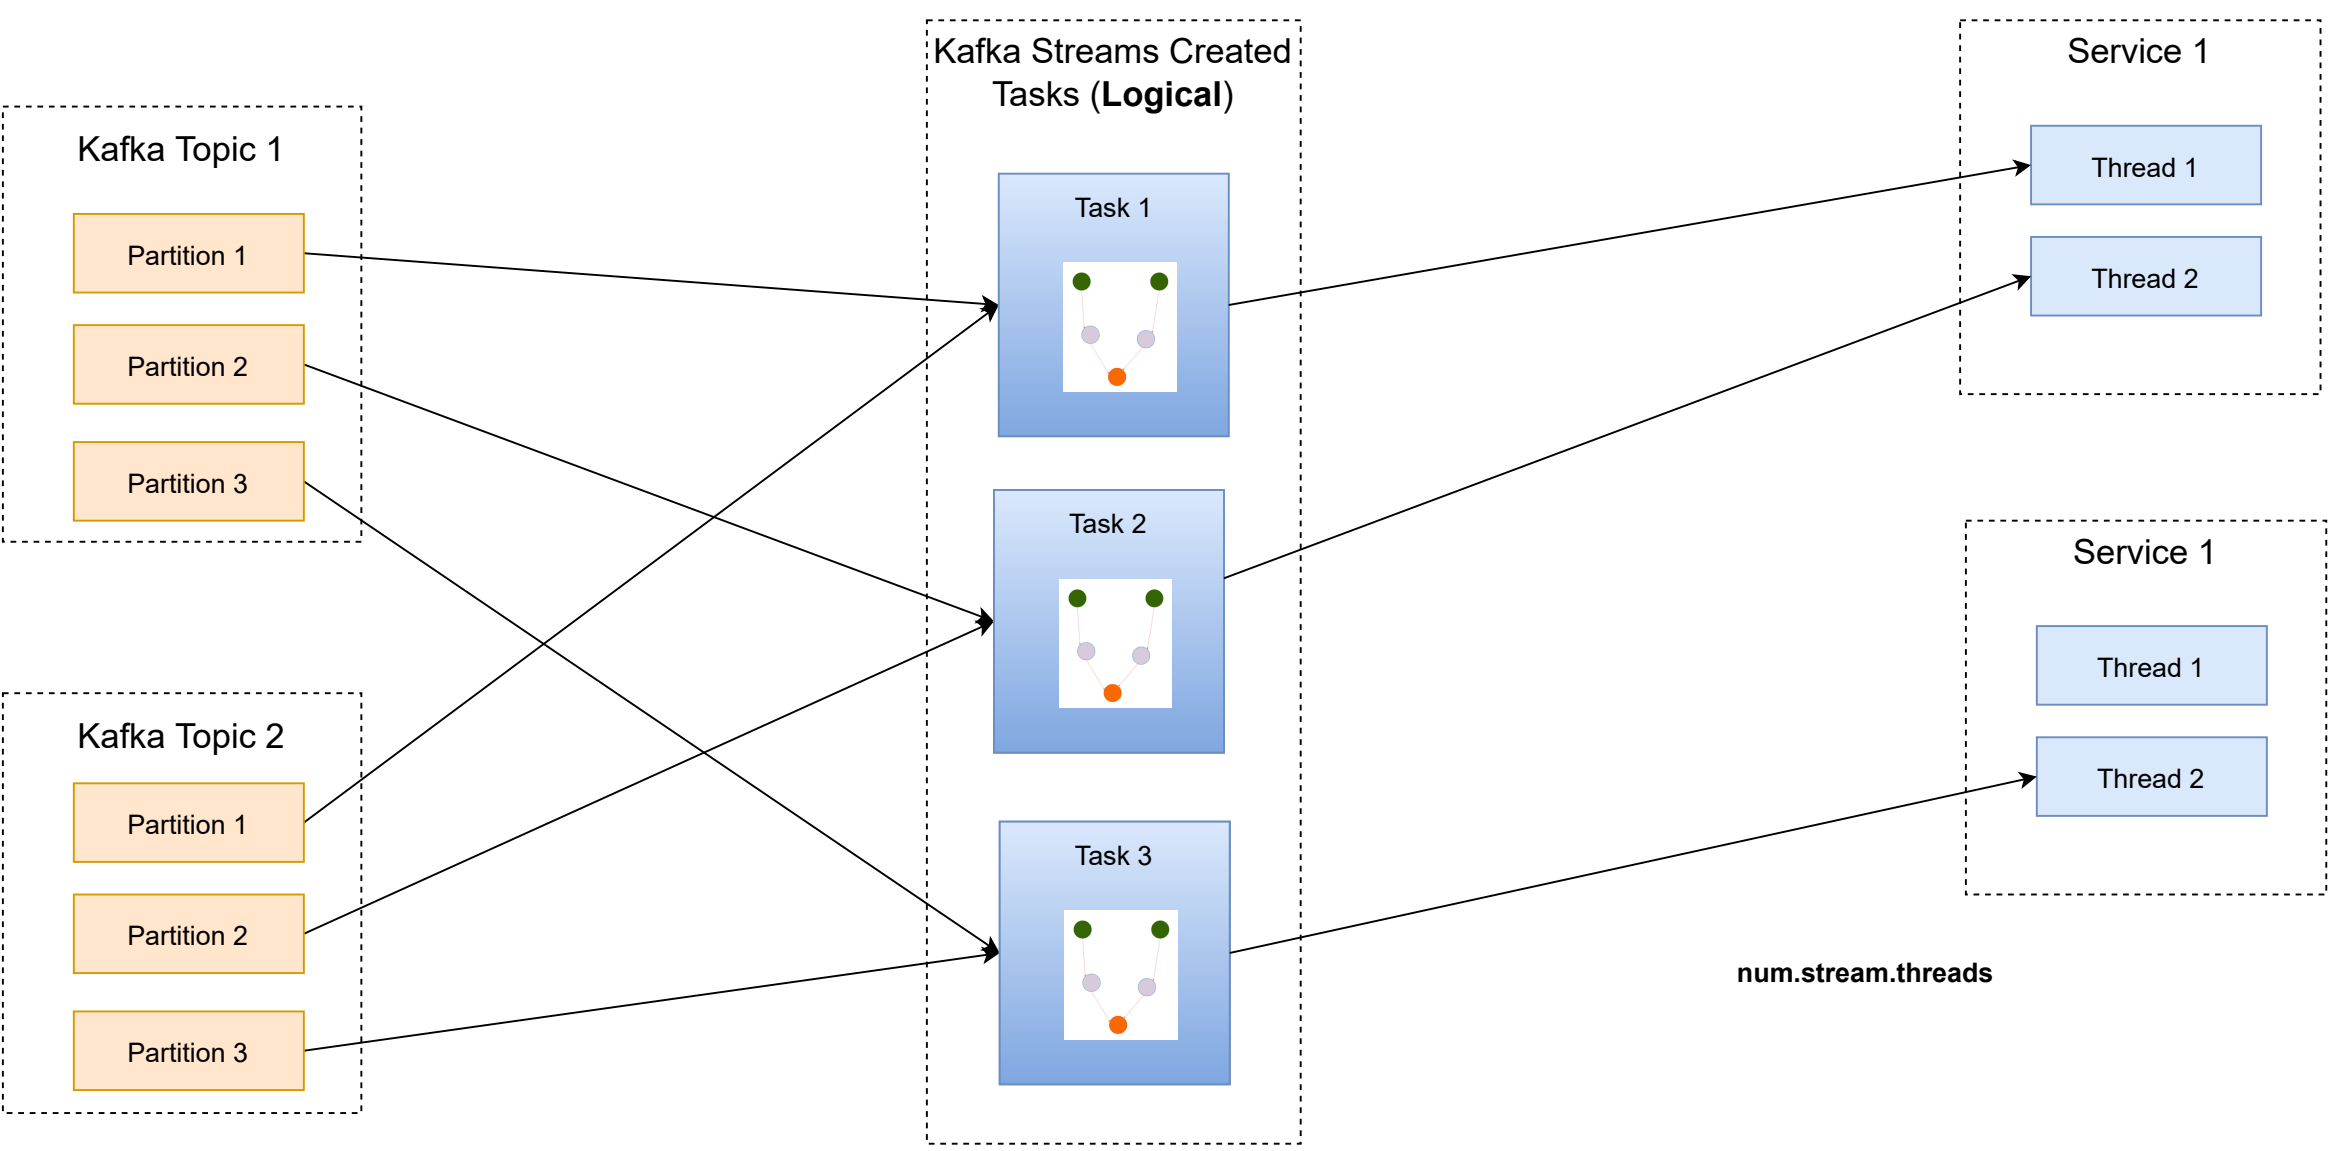


count



Write the favourite colors
with counts to Topic

Kafka Streaming Architecture



Input	KStream	KTable
Ashok:45000 Aravind:90000	Ashok:45000 Aravind:90000	Ashok:45000 Aravind:90000
Centhil:100000	Ashok:45000 Aravind:90000 Centhil:100000	Ashok:45000 Aravind:90000 Centhil:100000
Ashok:75000	Ashok:45000 Aravind:90000 Centhil:100000 Ashok:75000	Ashok:75000 Aravind:90000 Centhil:100000